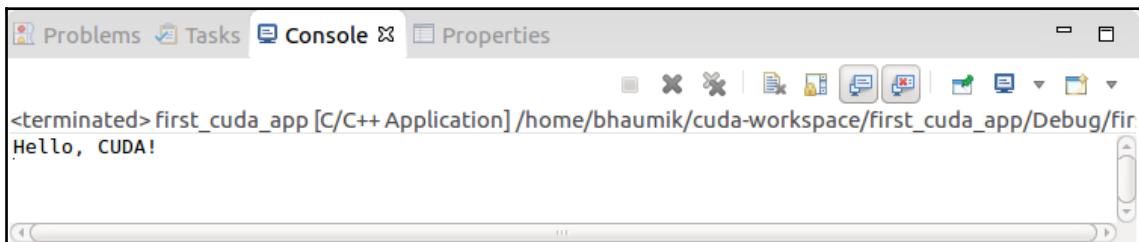


# Chapter 1: Introducing CUDA and Getting Started with CUDA

```
bhaumik@bhaumik-Lenovo-ideapad-520-15IKB: ~/NVIDIA_CUDA-9.0_Samples/1_Utilities/
Maximum number of threads per block: 1024
Max dimension size of a thread block (x,y,z): (1024, 1024, 64)
Max dimension size of a grid size (x,y,z): (2147483647, 65535, 65535)
Maximum memory pitch: 2147483647 bytes
Texture alignment: 512 bytes
Concurrent copy and kernel execution: Yes with 1 copy engine(s)
Run time limit on kernels: Yes
Integrated GPU sharing Host Memory: No
Support host page-locked memory mapping: Yes
Alignment requirement for Surfaces: Yes
Device has ECC support: Disabled
Device supports Unified Addressing (UVA): Yes
Supports Cooperative Kernel Launch: No
Supports MultiDevice Co-op Kernel Launch: .No
Device PCI Domain ID / Bus ID / location ID: 0 / 1 / 0
Compute Mode:
    < Default (multiple host threads can use ::cudaSetDevice() with device simultaneously) >

deviceQuery, CUDA Driver = CUDART, CUDA Driver Version = 9.0, CUDA Runtime Version = 9.0, NumDevs = 1
Result = PASS
bhaumik@bhaumik-Lenovo-ideapad-520-15IKB:~/NVIDIA_CUDA-9.0_Samples/1_Utilities/deviceQuery$
```



---

# Chapter 2: Parallel Programming using CUDA C

```
C:\WINDOWS\system32\cmd.exe  
1 + 4 = 5  
Press any key to continue . . .
```

```
C:\WINDOWS\system32\cmd.exe  
Passing Parameter by Reference Output: 1 + 4 = 5  
Press any key to continue . . .
```

```
C:\WINDOWS\system32\cmd.exe  
Hello!!!I'm thread in block: 2  
Hello!!!I'm thread in block: 7  
Hello!!!I'm thread in block: 6  
Hello!!!I'm thread in block: 8  
Hello!!!I'm thread in block: 1  
Hello!!!I'm thread in block: 0  
Hello!!!I'm thread in block: 5  
Hello!!!I'm thread in block: 10  
Hello!!!I'm thread in block: 9  
Hello!!!I'm thread in block: 11  
Hello!!!I'm thread in block: 4  
Hello!!!I'm thread in block: 3  
Hello!!!I'm thread in block: 14  
Hello!!!I'm thread in block: 13  
Hello!!!I'm thread in block: 12  
Hello!!!I'm thread in block: 15  
All threads are finished!  
Press any key to continue . . .
```

```
C:\WINDOWS\system32\cmd.exe
Detected 1 CUDA Capable device(s)

Device 0: "GeForce 940MX"
  CUDA Driver Version / Runtime Version          9.1 / 9.0
  CUDA Capability Major/Minor version number:    5.0
  Total amount of global memory:                4096 MBytes (4294967296 bytes)
  ( 3 ) Multiprocessors GPU Max Clock rate:      1189 MHz (1.19 GHz)
  Memory Clock rate:                          2505 Mhz
  Memory Bus Width:                           64-bit
  L2 Cache Size:                            1048576 bytes
  Maximum Texture Dimension Size (x,y,z)       1D=(65536), 2D=(65536, 65536), 3D=(4096, 4096, 4096)
  Maximum Layered 1D Texture Size, (num) layers   1D=(16384), 2048 layers
  Maximum Layered 2D Texture Size, (num) layers   2D=(16384, 16384), 2048 layers
  Total amount of constant memory:            65536 bytes
  Total amount of shared memory per block:     49152 bytes
  Total number of registers available per block: 65536
  Warp size:                                32
  Maximum number of threads per multiprocessor: 2048
  Maximum number of threads per block:        1024
  Max dimension size of a thread block (x,y,z): (1024, 1024, 64)
  Max dimension size of a grid size    (x,y,z): (2147483647, 65535, 65535)
  Maximum memory pitch:                      2147483647 bytes
  Texture alignment:                         512 bytes
  Concurrent copy and kernel execution:      Yes with 1 copy engine(s)
  Run time limit on kernels:                 Yes
  Integrated GPU sharing Host Memory:        No
  Support host page-locked memory mapping:   Yes
  Alignment requirement for Surfaces:        Yes
  Device has ECC support:                   Disabled
  CUDA Device Driver Mode (TCC or WDDM):    WDDM (Windows Display Driver Model)
  Device supports Unified Addressing (UVA): Yes
  Supports Cooperative Kernel Launch:       No
  Supports MultiDevice Co-op Kernel Launch: No
  Device PCI Domain ID / Bus ID / location ID: 0 / 1 / 0
  Compute Mode:
    < Default (multiple host threads can use ::cudaSetDevice() with device simultaneously) >
Press any key to continue . . .
```

```
C:\WINDOWS\system32\cmd.exe
Vector addition on CPU
The sum of 0 element is 0 + 0 = 0
The sum of 1 element is 2 + 1 = 3
The sum of 2 element is 8 + 2 = 10
The sum of 3 element is 18 + 3 = 21
The sum of 4 element is 32 + 4 = 36
Press any key to continue . . .
```

---

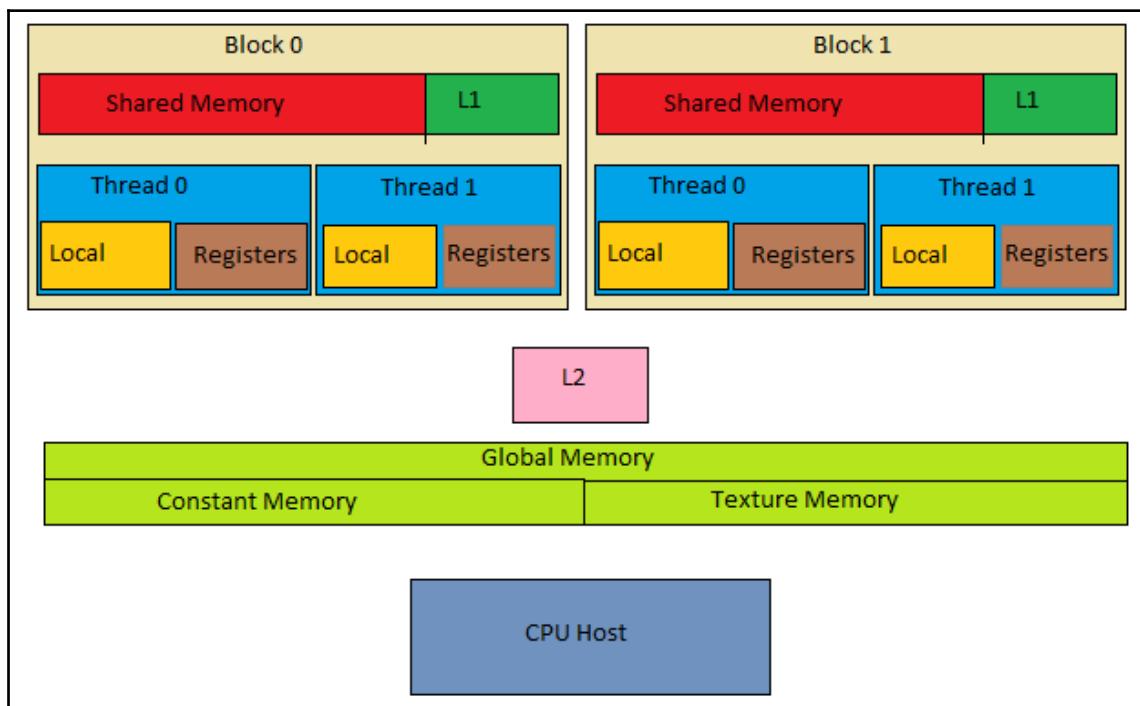
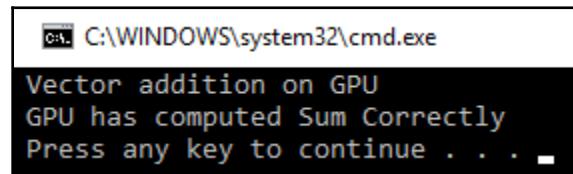
```
C:\ C:\WINDOWS\system32\cmd.exe
Vector addition on GPU
The sum of 0 element is 0 + 0 = 0
The sum of 1 element is 2 + 1 = 3
The sum of 2 element is 8 + 2 = 10
The sum of 3 element is 18 + 3 = 21
The sum of 4 element is 32 + 4 = 36
Press any key to continue . . . ■
```

```
No of Elements in Array:10000000
Device time 0.001000 seconds
host time 0.025000 Seconds
Press any key to continue . . .
```

```
C:\ C:\WINDOWS\system32\cmd.exe
Square of Number on GPU
The square of 0.000000 is 0.000000
The square of 1.000000 is 1.000000
The square of 2.000000 is 4.000000
The square of 3.000000 is 9.000000
The square of 4.000000 is 16.000000
Press any key to continue . . . ■
```

# Chapter 3: Threads, Synchronization, and Memory

Block 0	Thread 0	Thread 1	Thread 2
Block 1	Thread 0	Thread 1	Thread 2
Block 2	Thread 0	Thread 1	Thread 2



---

```
C:\WINDOWS\system32\cmd.exe
Array in Global Memory is:
At Index: 0 --> 0
At Index: 1 --> 1
At Index: 2 --> 2
At Index: 3 --> 3
At Index: 4 --> 4
Press any key to continue . . .
```

```
C:\WINDOWS\system32\cmd.exe
Use of Local Memory on GPU:
Value of Local variable in current thread is: 0
Value of Local variable in current thread is: 5
Value of Local variable in current thread is: 10
Value of Local variable in current thread is: 15
Value of Local variable in current thread is: 20
Press any key to continue . . .
```

```
C:\WINDOWS\system32\cmd.exe
Use of Shared Memory on GPU:
The running average after 0 element is 0.000000
The running average after 1 element is 0.500000
The running average after 2 element is 1.000000
The running average after 3 element is 1.500000
The running average after 4 element is 2.000000
The running average after 5 element is 2.500000
The running average after 6 element is 3.000000
The running average after 7 element is 3.500000
The running average after 8 element is 4.000000
The running average after 9 element is 4.500000
Press any key to continue . . .
```

```
C:\WINDOWS\system32\cmd.exe
10000 total threads in 100 blocks writing into 10 array elements
Number of times a particular Array index has been incremented without atomic add is:
index: 0 --> 16 times
index: 1 --> 16 times
index: 2 --> 16 times
index: 3 --> 16 times
index: 4 --> 16 times
index: 5 --> 16 times
index: 6 --> 17 times
index: 7 --> 17 times
index: 8 --> 17 times
index: 9 --> 17 times
Press any key to continue . . .
```

```
C:\WINDOWS\system32\cmd.exe
10000 total threads in 100 blocks writing into 10 array elements
Number of times a particular Array index has been incremented without atomic add is:
index: 0 --> 19 times
index: 1 --> 19 times
index: 2 --> 19 times
index: 3 --> 19 times
index: 4 --> 19 times
index: 5 --> 19 times
index: 6 --> 19 times
index: 7 --> 19 times
index: 8 --> 19 times
index: 9 --> 19 times
Press any key to continue . . .
```

```
C:\WINDOWS\system32\cmd.exe
10000 total threads in 100 blocks writing into 10 array elements
Number of times a particular Array index has been incremented is:
index: 0 --> 1000 times
index: 1 --> 1000 times
index: 2 --> 1000 times
index: 3 --> 1000 times
index: 4 --> 1000 times
index: 5 --> 1000 times
index: 6 --> 1000 times
index: 7 --> 1000 times
index: 8 --> 1000 times
index: 9 --> 1000 times
Press any key to continue . . .
```

```
C:\WINDOWS\system32\cmd.exe
Use of Constant memory on GPU
The expression for input 0.000000 is 20.000000
The expression for input 1.000000 is 22.000000
The expression for input 2.000000 is 24.000000
The expression for input 3.000000 is 26.000000
The expression for input 4.000000 is 28.000000
Press any key to continue . . .
```

```
C:\WINDOWS\system32\cmd.exe
Use of Texture memory on GPU:
Texture element at 0 is : 0.000000
Texture element at 1 is : 1.000000
Texture element at 2 is : 2.000000
Texture element at 3 is : 3.000000
Texture element at 4 is : 4.000000
Texture element at 5 is : 5.000000
Texture element at 6 is : 6.000000
Texture element at 7 is : 7.000000
Texture element at 8 is : 8.000000
Texture element at 9 is : 9.000000
Press any key to continue . . .
```

```
C:\WINDOWS\system32\cmd.exe
The computed dot product is: 1047552.000000
The dot product computed by GPU is correct
Press any key to continue . . .
```

$$\begin{bmatrix} 0 & 0 & 0 & 0 \\ 1 & 1 & 1 & 1 \\ 2 & 2 & 2 & 2 \\ 3 & 3 & 3 & 3 \end{bmatrix} * \begin{bmatrix} 0 & 1 & 2 & 3 \\ 0 & 1 & 2 & 3 \\ 0 & 1 & 2 & 3 \\ 0 & 1 & 2 & 3 \end{bmatrix} = \begin{bmatrix} 0 * 0 + 0 * 0 + 0 * 0 + 0 * 0 & 0 & 0 & 0 \\ 1 * 0 + 1 * 0 + 1 * 0 + 1 * 0 & 4 & 8 & 12 \\ 2 * 0 + 2 * 0 + 2 * 0 + 2 * 0 & 8 & 16 & 24 \\ 3 * 0 + 3 * 0 + 3 * 0 + 3 * 0 & 12 & 24 & 36 \end{bmatrix}$$

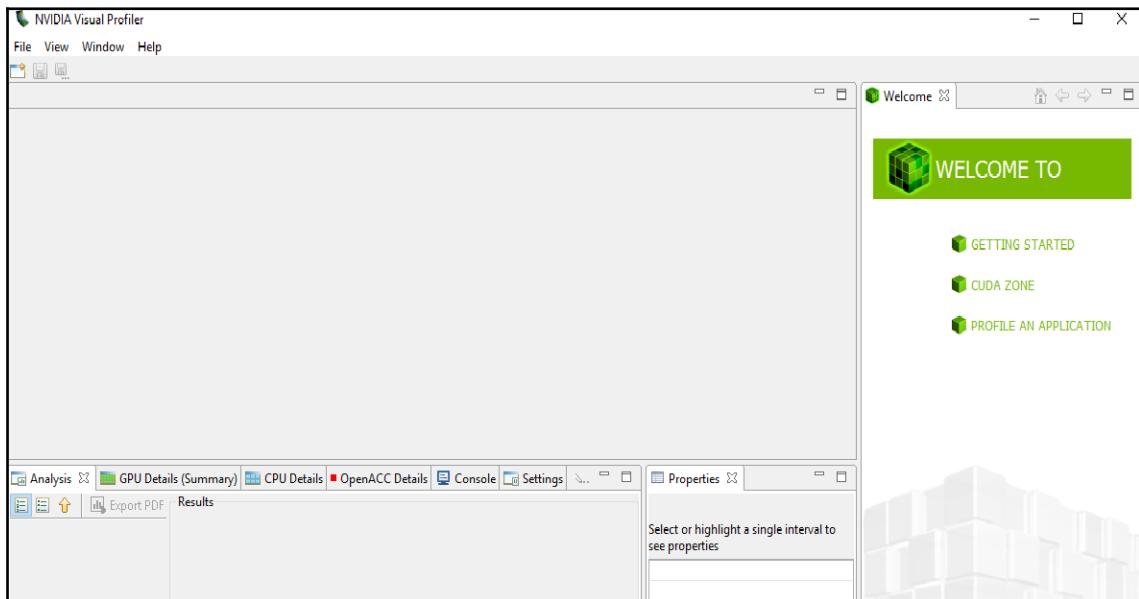
$$\begin{bmatrix} M_{00} & M_{01} \\ M_{10} & M_{11} \end{bmatrix} \rightarrow [M_{00} \quad M_{01} \quad M_{10} \quad M_{11}]$$

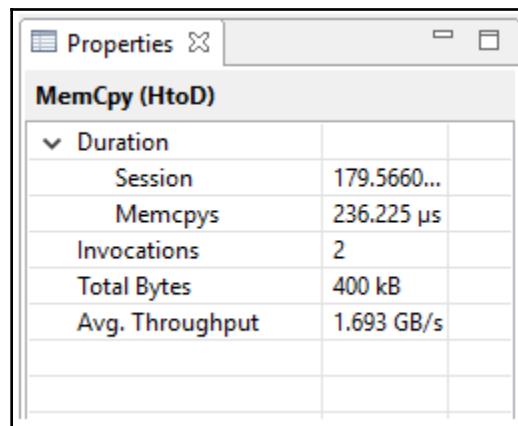
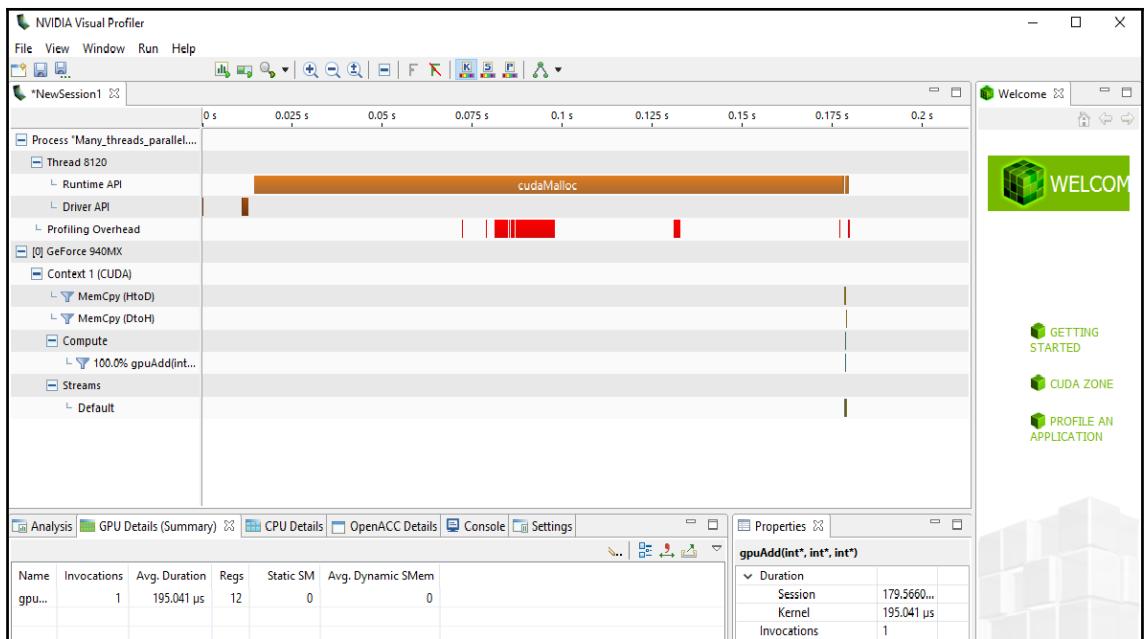
C:\ C:\WINDOWS\system32\cmd.exe

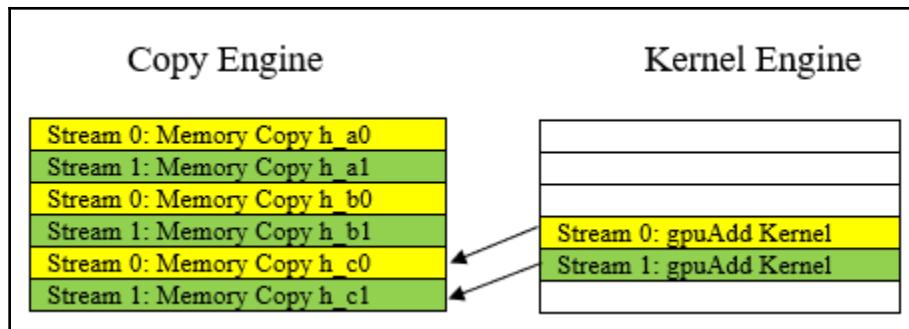
```
The result of Matrix multiplication is:  
0.000000 0.000000 0.000000 0.000000 0.000000 0.000000  
0.000000 6.000000 12.000000 18.000000 24.000000 30.000000  
0.000000 12.000000 24.000000 36.000000 48.000000 60.000000  
0.000000 18.000000 36.000000 54.000000 72.000000 90.000000  
0.000000 24.000000 48.000000 72.000000 96.000000 120.000000  
0.000000 30.000000 60.000000 90.000000 120.000000 150.000000  
Press any key to continue . . . -
```

# Chapter 4: Advanced Concepts in CUDA

```
C:\WINDOWS\system32\cmd.exe
Time to add 50000 numbers: 0.9 ms
Vector addition on GPU
GPU has computed Sum Correctly
Press any key to continue . . .
```







```
c:\ C:\WINDOWS\system32\cmd.exe
Time to add 100000 numbers: 0.9 ms
Vector addition on GPU
GPU has computed Sum Correctly
Press any key to continue . . . -
```

```
c:\ C:\WINDOWS\system32\cmd.exe
The Enumeration sorted Array is:
3
4
5
8
9
Press any key to continue . . .
```

---

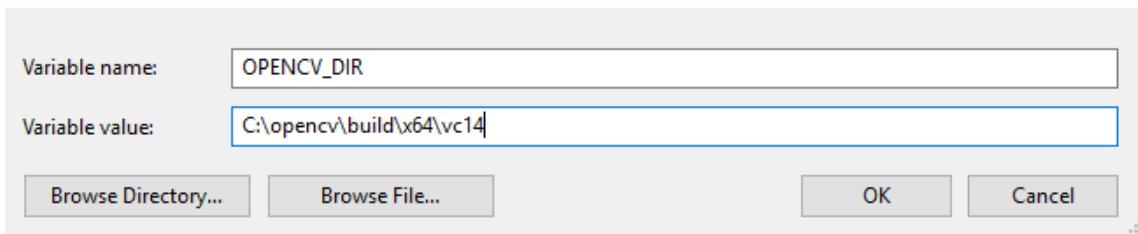
```
C:\> C:\WINDOWS\system32\cmd.exe
Histogram using 16 bin without shared Memory is:
bin 0: count 63
bin 1: count 63
bin 2: count 63
bin 3: count 63
bin 4: count 63
bin 5: count 63
bin 6: count 63
bin 7: count 63
bin 8: count 62
bin 9: count 62
bin 10: count 62
bin 11: count 62
bin 12: count 62
bin 13: count 62
bin 14: count 62
bin 15: count 62
Press any key to continue . . .
```

```
C:\> C:\WINDOWS\system32\cmd.exe
bin 227: count 4
bin 228: count 4
bin 229: count 4
bin 230: count 4
bin 231: count 4
bin 232: count 3
bin 233: count 3
bin 234: count 3
bin 235: count 3
bin 236: count 3
bin 237: count 3
bin 238: count 3
bin 239: count 3
bin 240: count 3
bin 241: count 3
bin 242: count 3
bin 243: count 3
bin 244: count 3
bin 245: count 3
bin 246: count 3
bin 247: count 3
bin 248: count 3
bin 249: count 3
bin 250: count 3
bin 251: count 3
bin 252: count 3
bin 253: count 3
bin 254: count 3
bin 255: count 3
Press any key to continue . . .-
```

# Chapter 5: Getting Started with OpenCV with CUDA Support

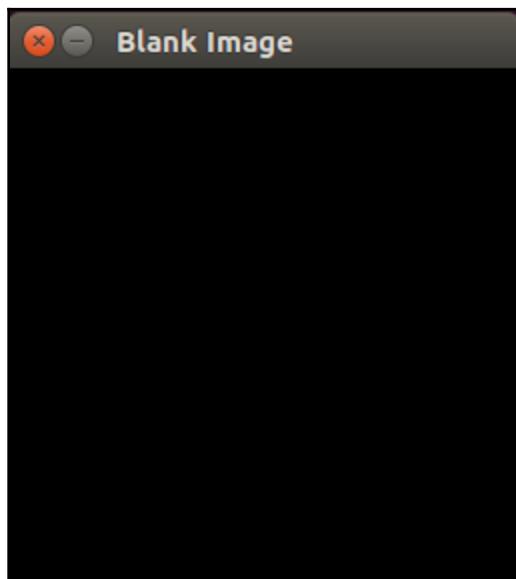
New System Variable

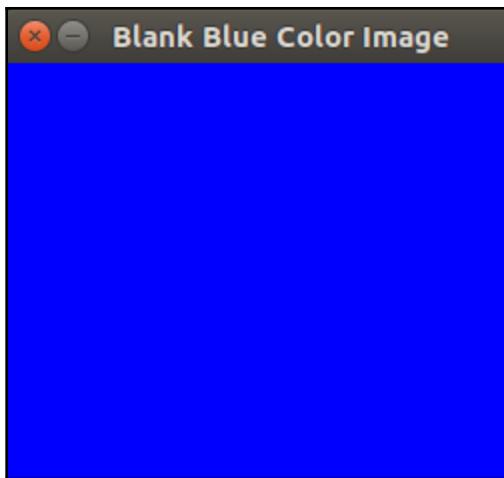
X

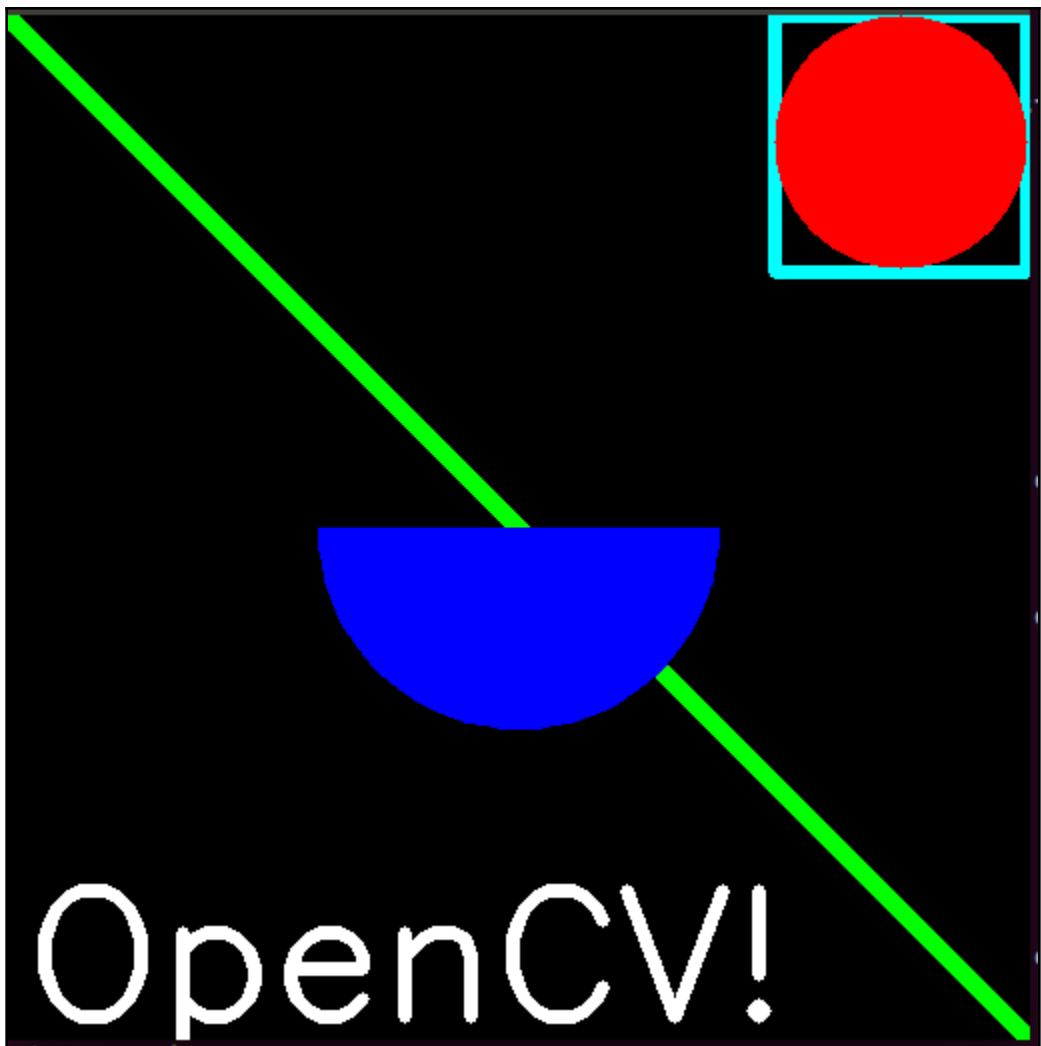


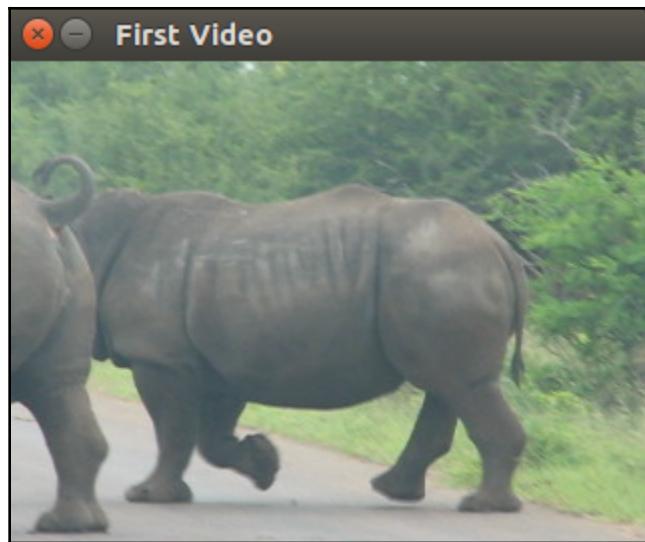
```
bhaumik@bhaumik-Lenovo-ideapad-520-15IKB: ~/Desktop/opencv/opencv/build
-- Parallel framework: pthreads
-- Trace: YES (with Intel ITT)
-- Other third-party libraries:
--   Lapack: NO
--   Eigen: NO
--   Custom HAL: NO
--   Protobuf: build (3.5.1)
-- NVIDIA CUDA: YES (ver 7.5, CUFFT CUBLAS FAST_MATH)
--   NVIDIA GPU archs: 20 30 35 37 50 52 60 61
--   NVIDIA PTX archs:
-- OpenCL:
--   Include path: /home/bhaumik/Desktop/opencv/opencv/3rdparty/include/opencl/1.2
--   Link libraries: Dynamic load
-- Python 2:
--   Interpreter: /usr/bin/python2.7 (ver 2.7.12)
--   Libraries: /usr/lib/x86_64-linux-gnu/libpython2.7.so (ver 2.7.12)
--   numpy: /usr/lib/python2.7/dist-packages/numpy/core/include (ver 1.11.0)
--   packages path: lib/python2.7/dist-packages
-- Python (for build):
--   Pylint: /usr/bin/pylint (ver: 1.8.2, checks: 147)
-- Java:
--   ant: NO
--   JNI: NO
--   Java wrappers: NO
--   Java tests: NO
-- Matlab: NO
-- Install to: /usr/local
-- Configuring done
-- Generating done
-- Build files have been written to: /home/bhaumik/Desktop/opencv/opencv/build
bhaumik@bhaumik-Lenovo-ideapad-520-15IKB:~/Desktop/opencv/opencv/build$
```

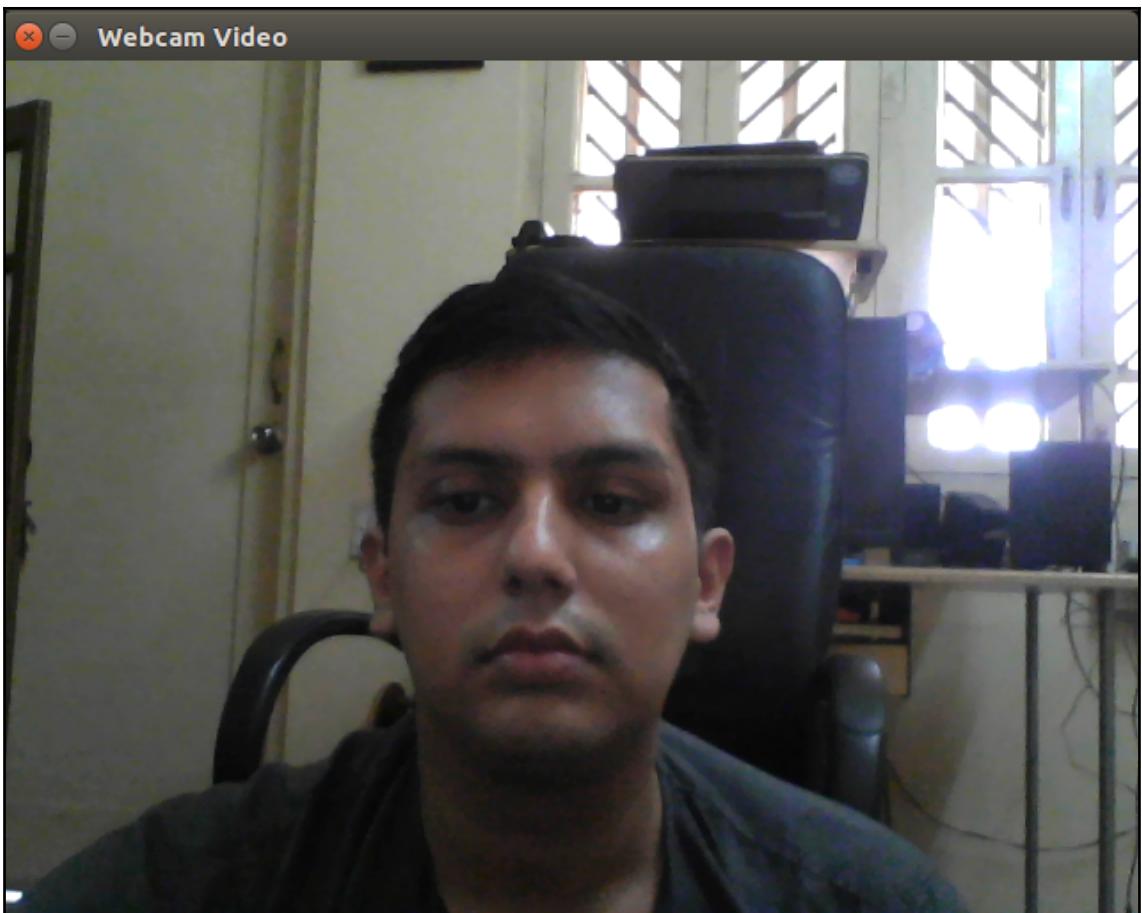


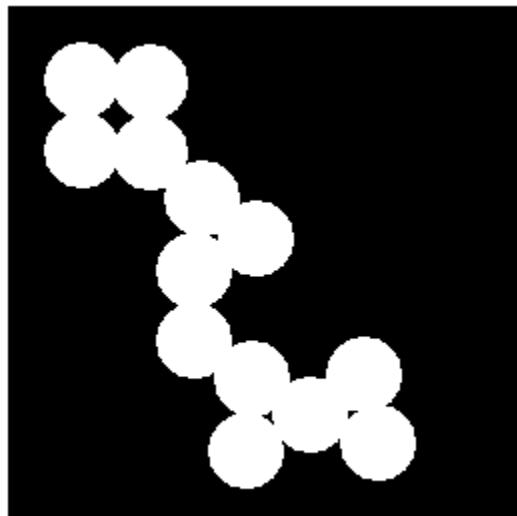




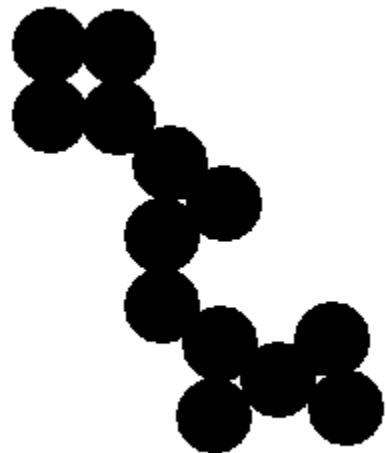








Source Image



Inverted Image

Performance of Thresholding on CPU:

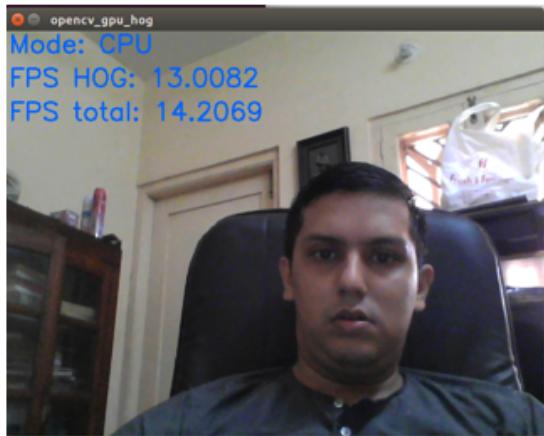
Time: 0.169766

FPS: 5.89046

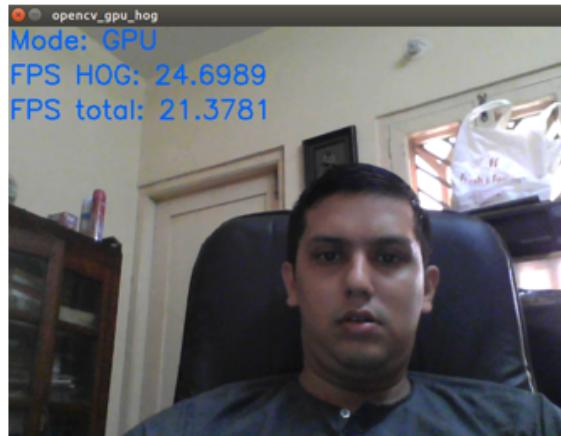
Performance of Thresholding on GPU:

Time: 0.000550593

FPS: 1816.22



CPU Performance

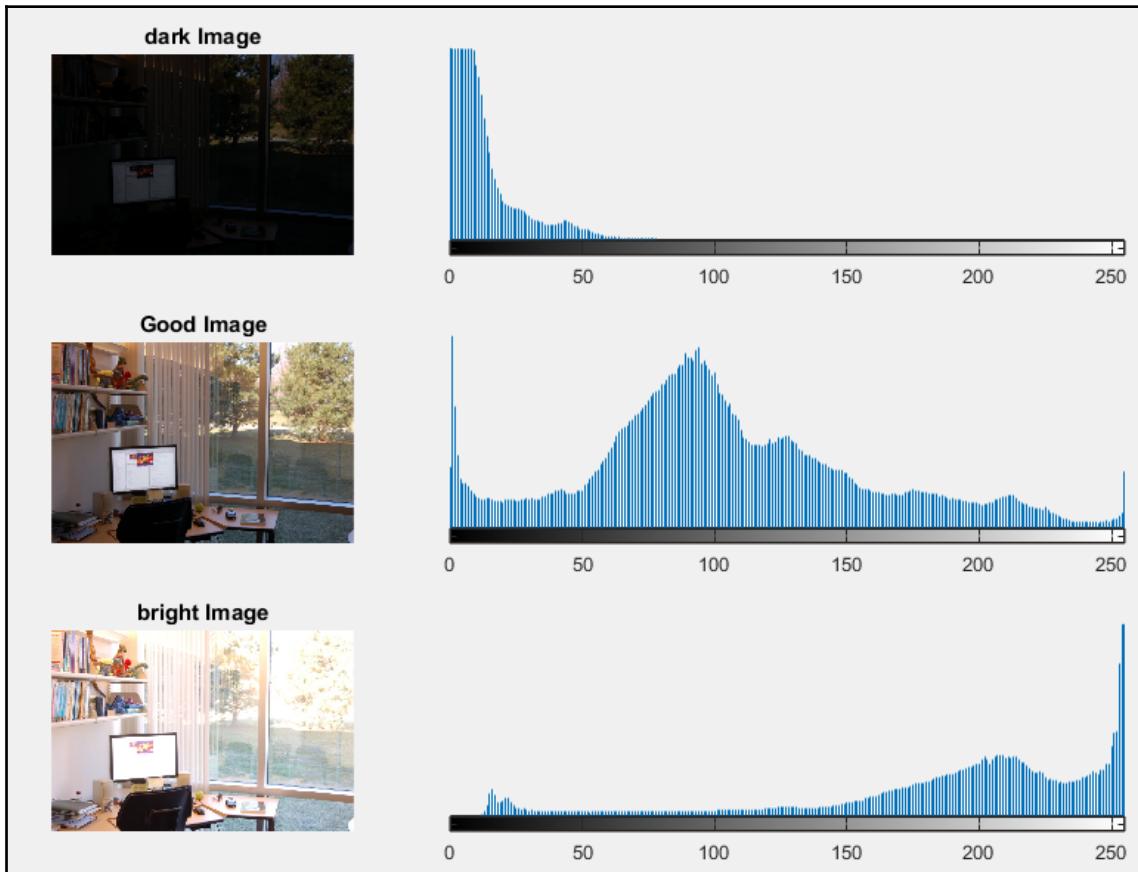


GPU Performance

---

# Chapter 6: Basic Computer Vision Operations Using OpenCV and CUDA

```
bhaumik@bhaumik-Lenovo-ideapad-520-15IKB:~/Desktop/opencv/Chapter 6$ g++ -std=c++11 individual_pixel.cpp `pkg-config --libs --cflags opencv` -o pixel  
bhaumik@bhaumik-Lenovo-ideapad-520-15IKB:~/Desktop/opencv/Chapter 6$ ./pixel  
Pixel Intensity of gray scale Image at (100,50) is:9  
Pixel Intensity of color Image at (100,50) is:[175, 179, 177]
```





---

Original Image



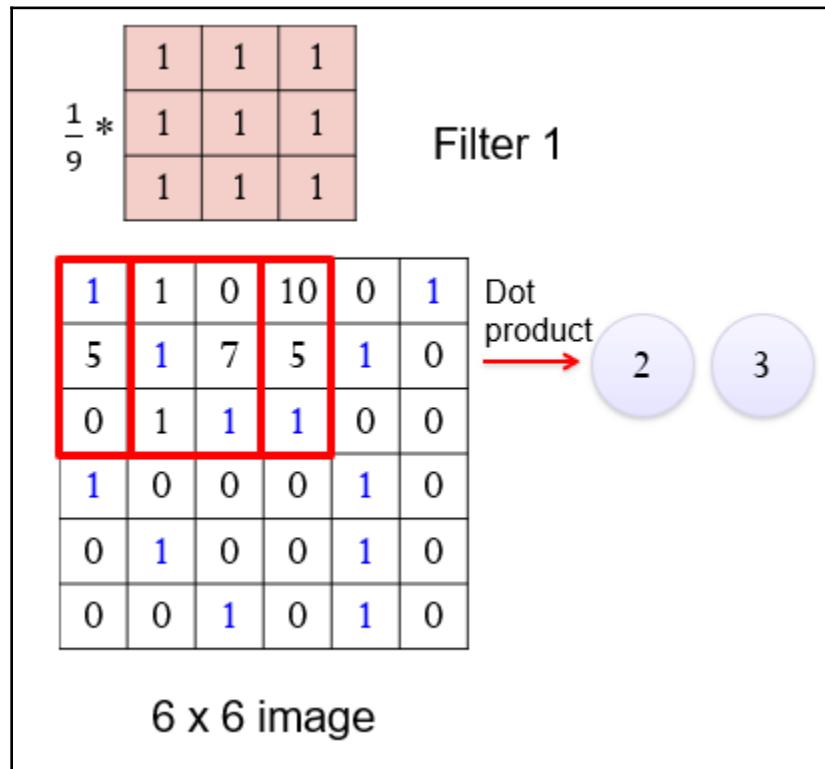
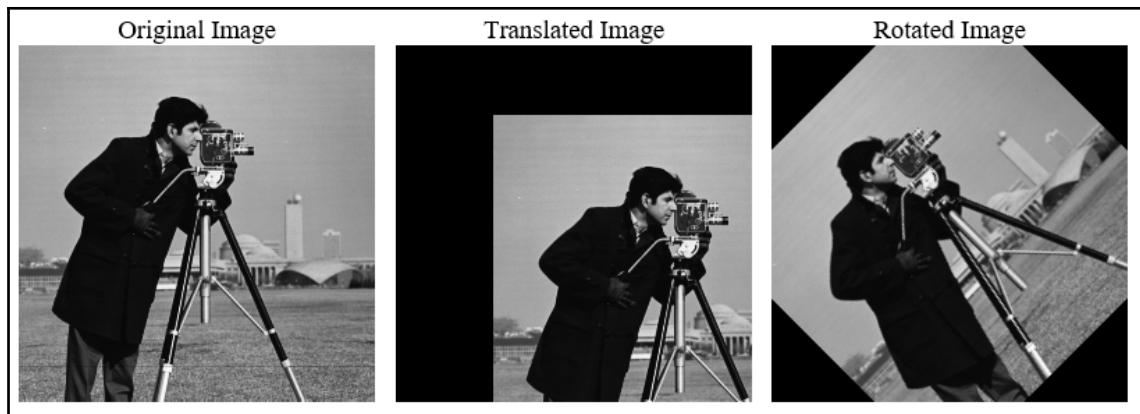
Image Resize (200,200)



Image Resize  
downscaled by 2



$$\begin{bmatrix} 1 & 0 & tx \\ 0 & 1 & ty \end{bmatrix}$$



---

Averaging Filter 3x3

$$\frac{1}{9} * \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

Averaging Filter 5x5

$$\frac{1}{25} * \begin{bmatrix} 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \end{bmatrix}$$

Original Image



Average Filter (3x3)



Average Filter (5x5)



Average Filter (7x7)



---

Gaussian Filter 5x5

$$\frac{1}{273} * \begin{bmatrix} 1 & 4 & 7 & 4 & 1 \\ 4 & 16 & 26 & 16 & 4 \\ 7 & 26 & 41 & 26 & 7 \\ 4 & 16 & 26 & 16 & 4 \\ 1 & 4 & 7 & 4 & 1 \end{bmatrix}$$

Original Image



Gaussian Filter (3x3)



Gaussian Filter (5x5)



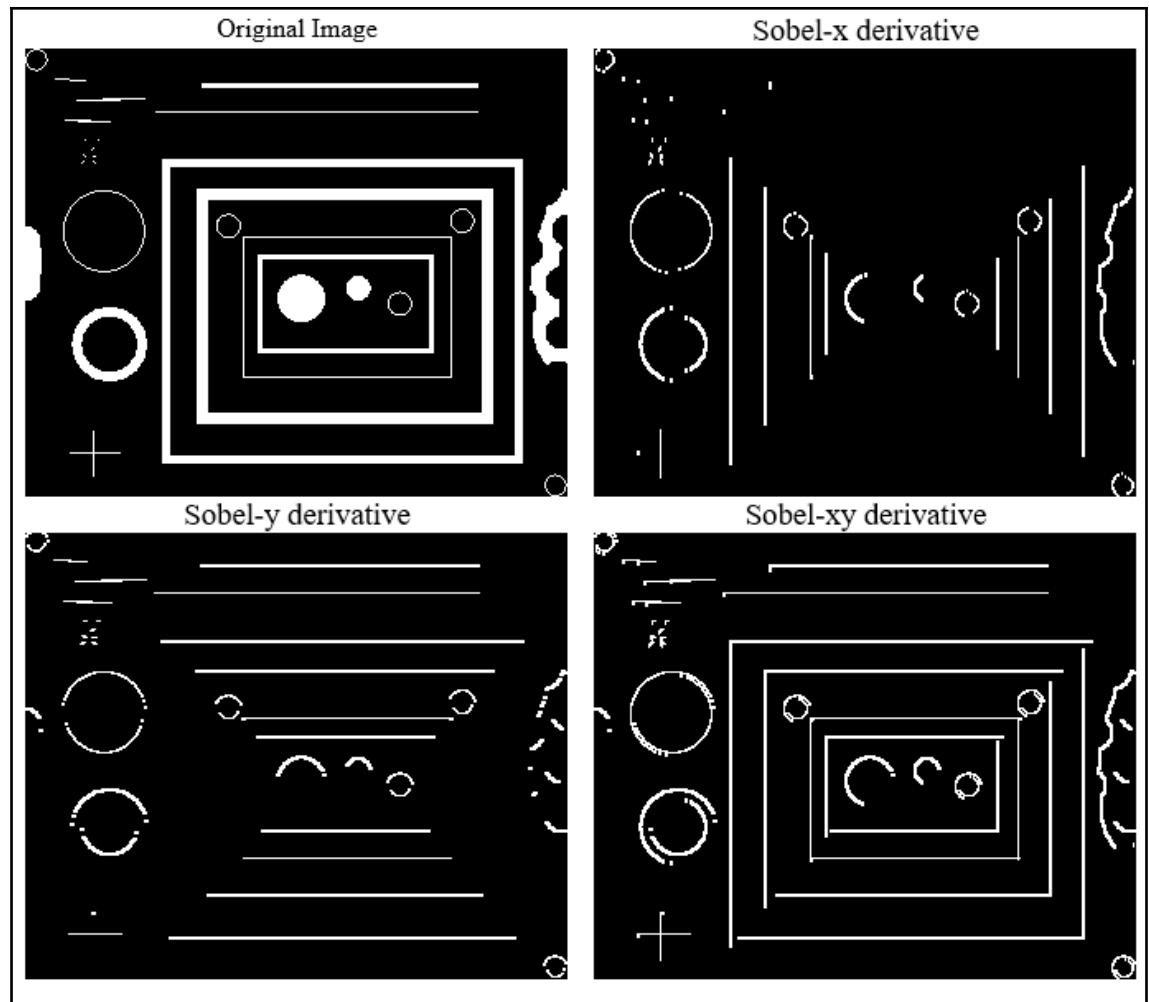
Gaussian Filter (7x7)





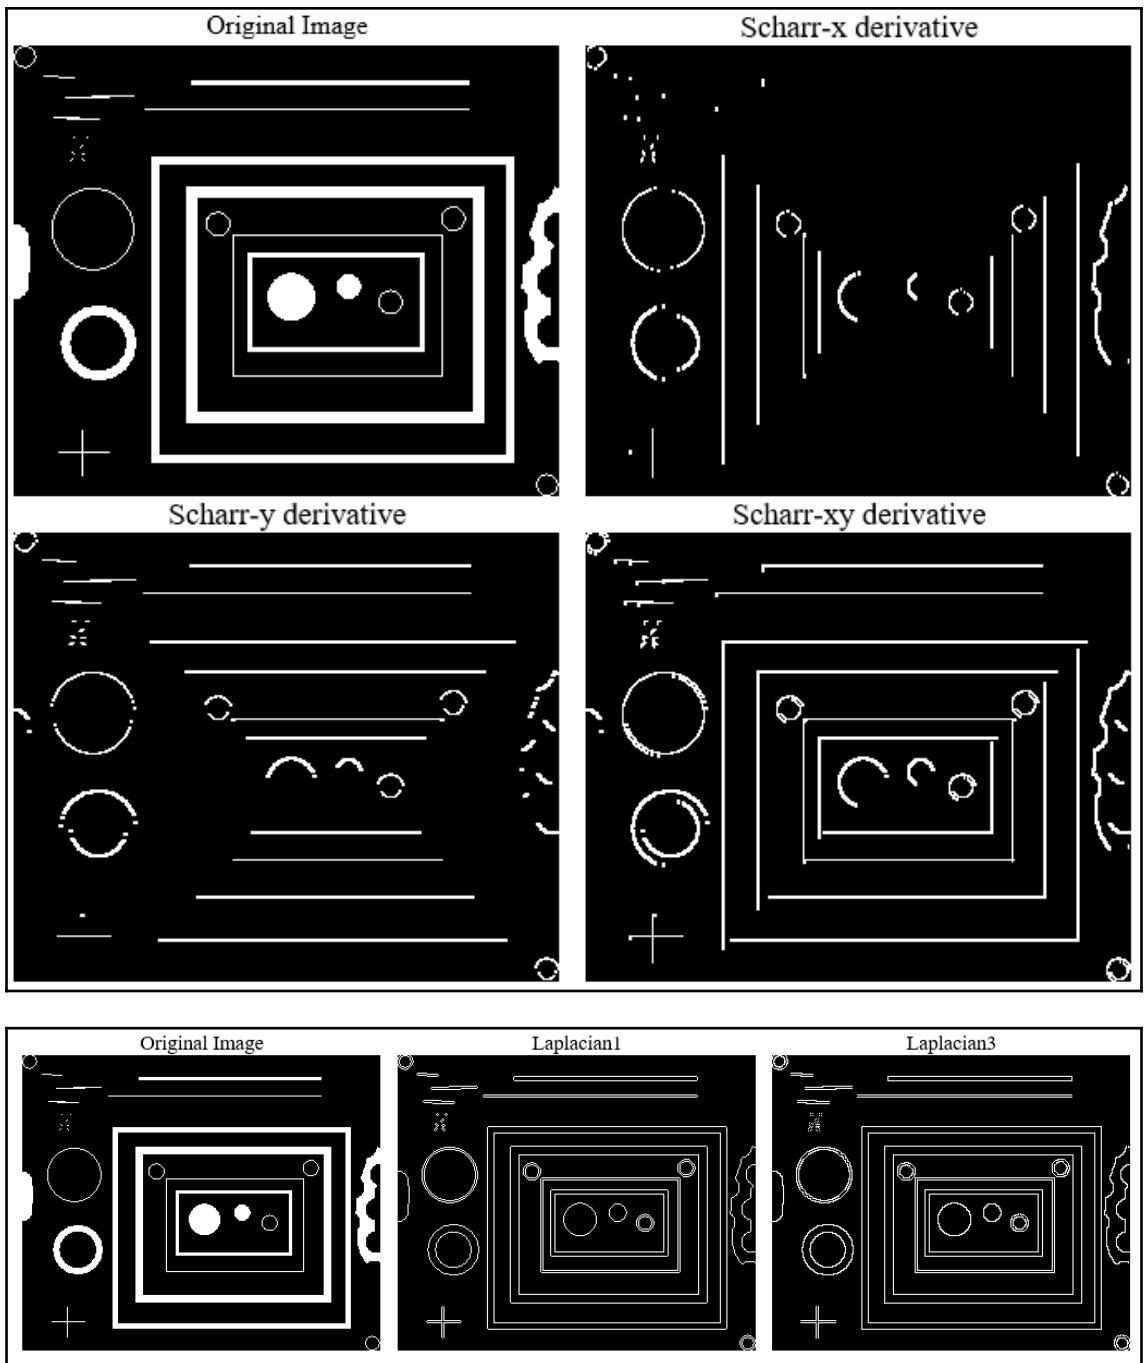
$$Sx = \begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}$$

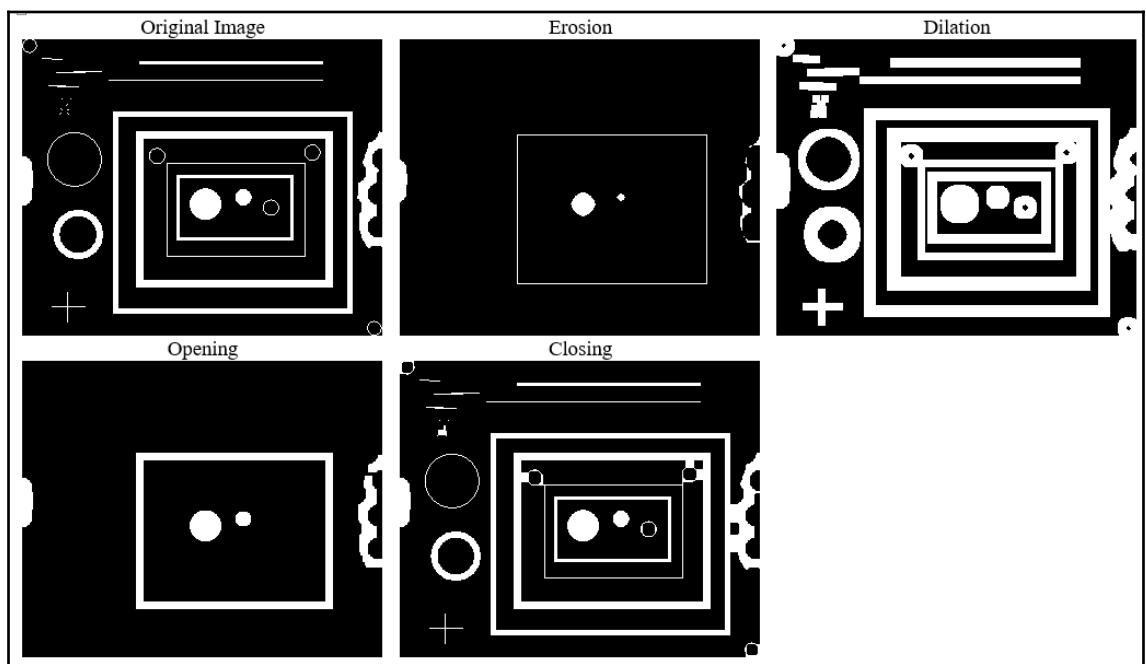
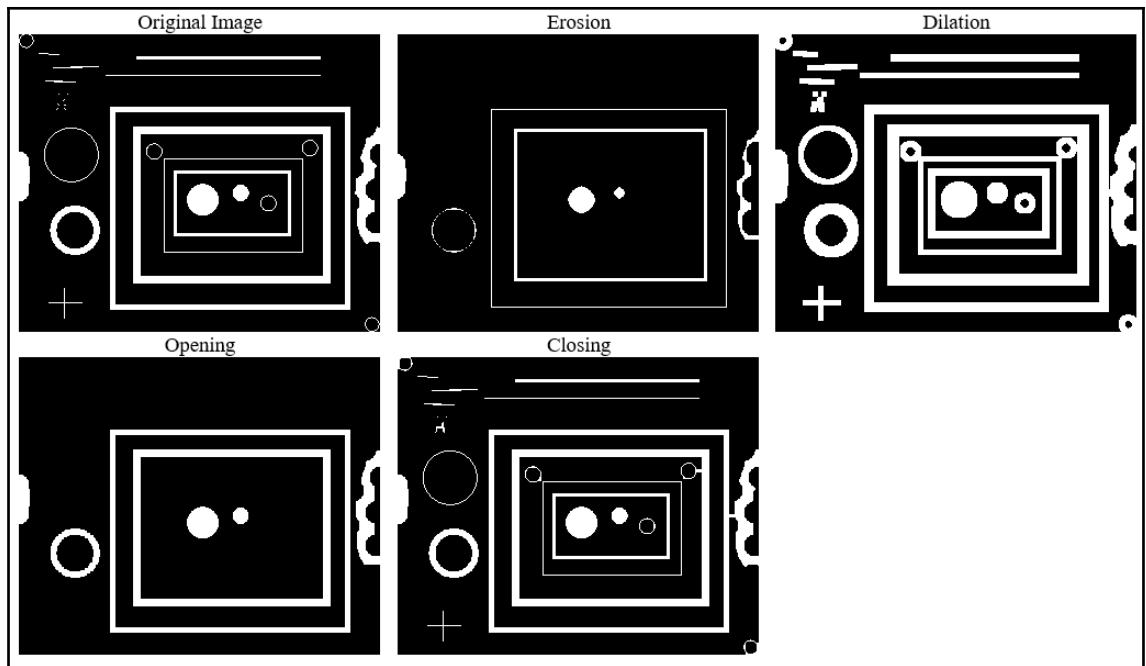
$$Sy = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$



$$Sx = \begin{bmatrix} 3 & 0 & -3 \\ 10 & 0 & -10 \\ 3 & 0 & -3 \end{bmatrix}$$

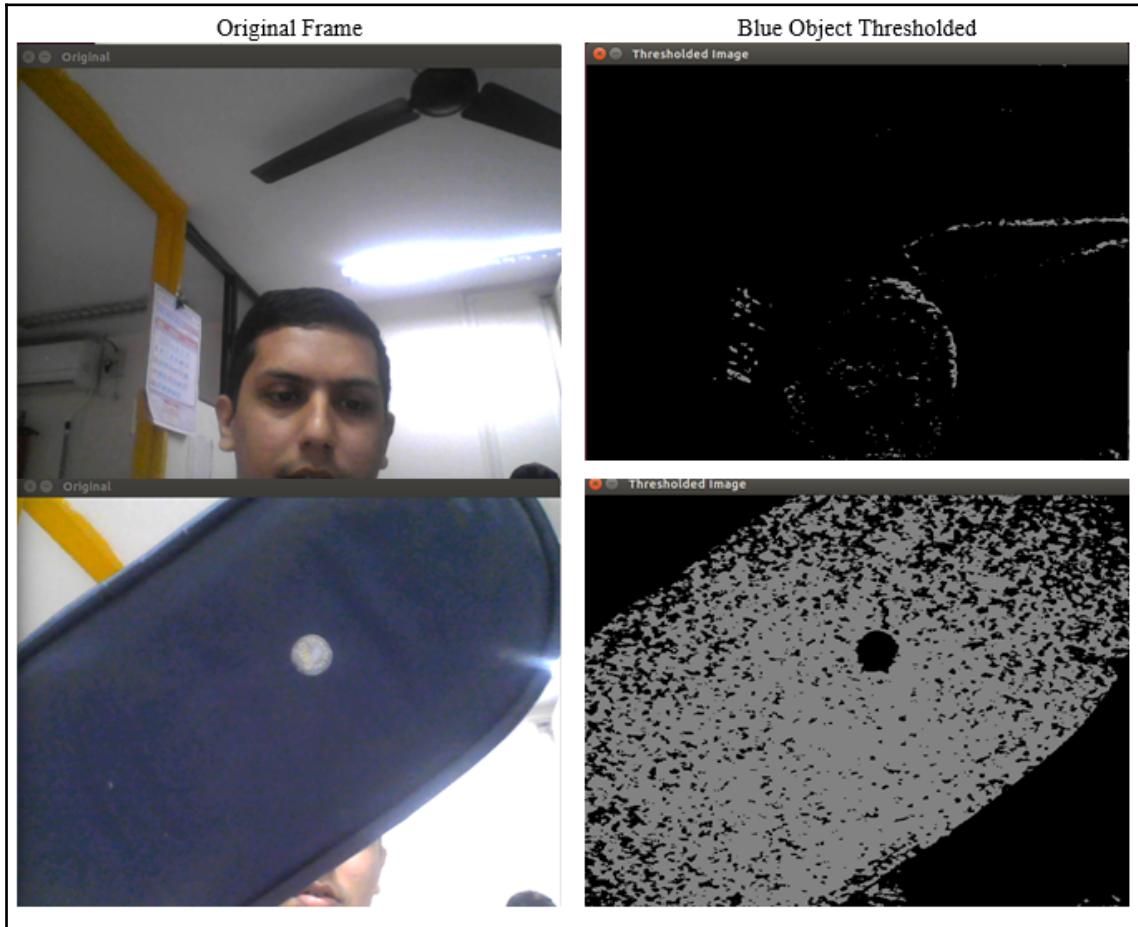
$$Sy = \begin{bmatrix} 3 & 10 & 3 \\ 0 & 0 & 0 \\ -3 & -10 & -3 \end{bmatrix}$$

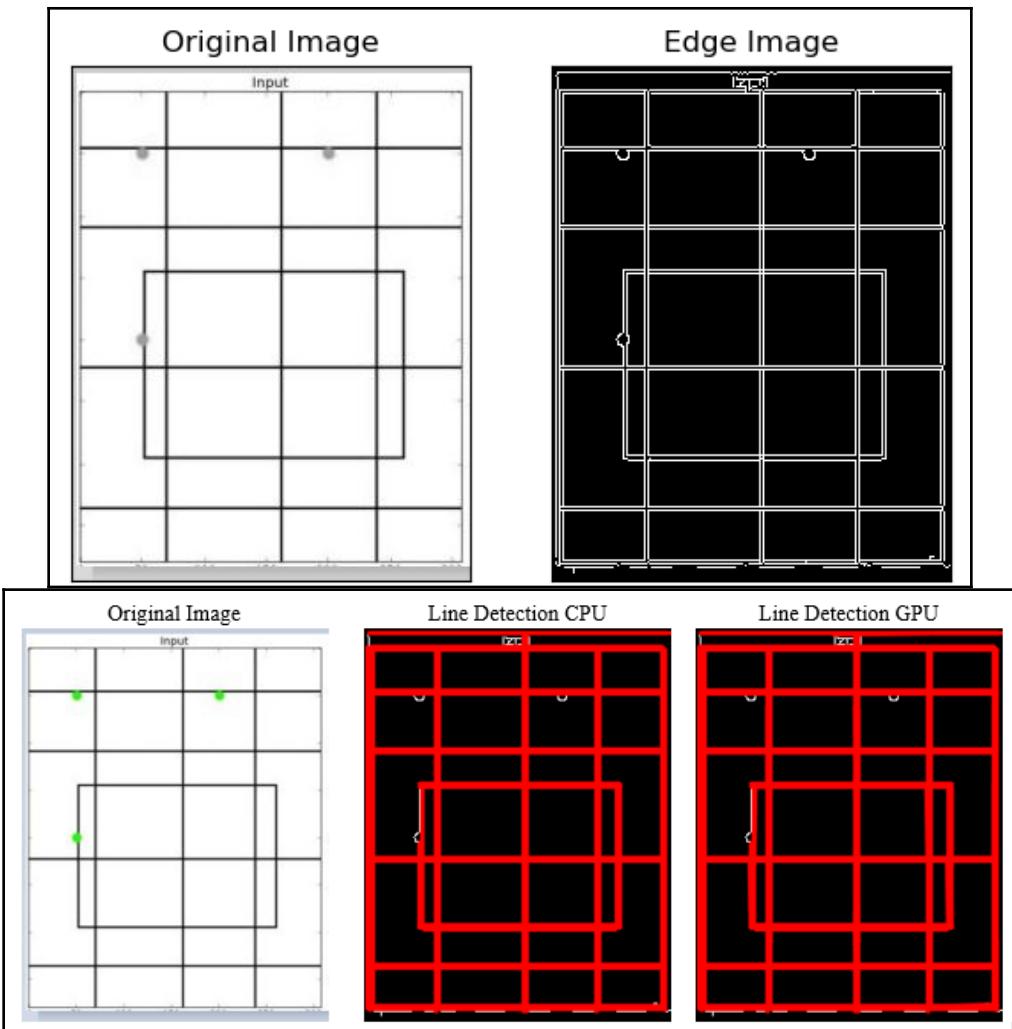




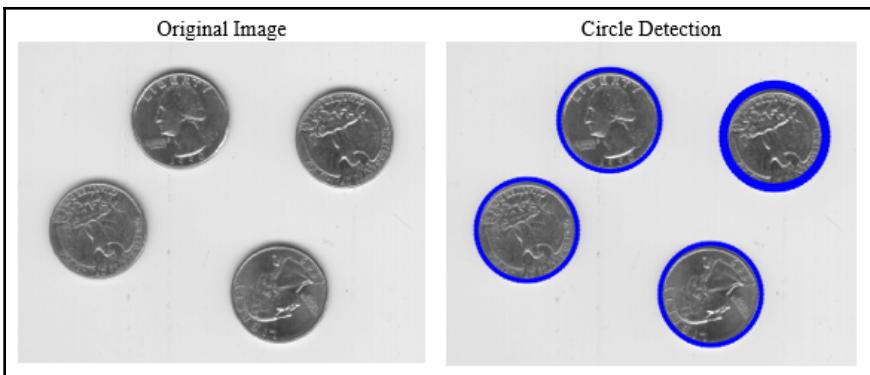
---

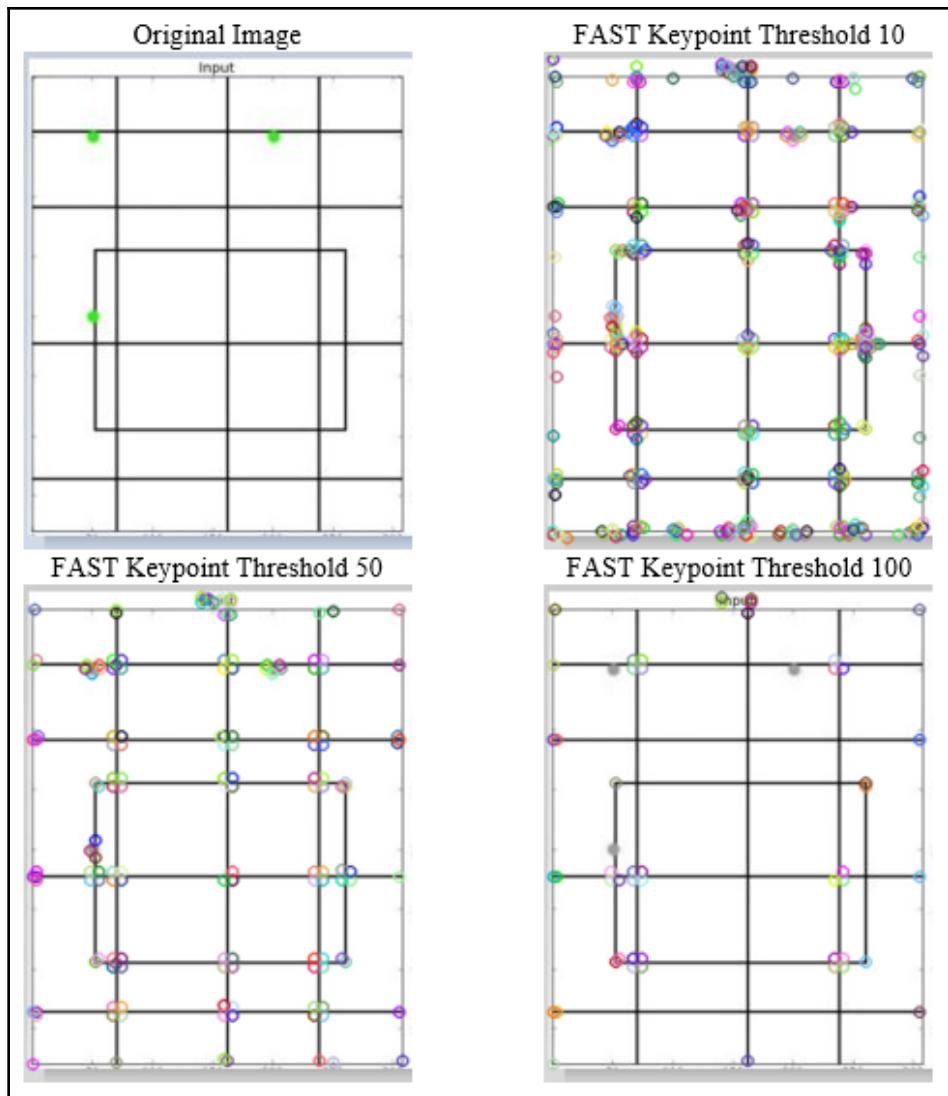
# Chapter 7: Object Detection and Tracking Using OpenCV and CUDA

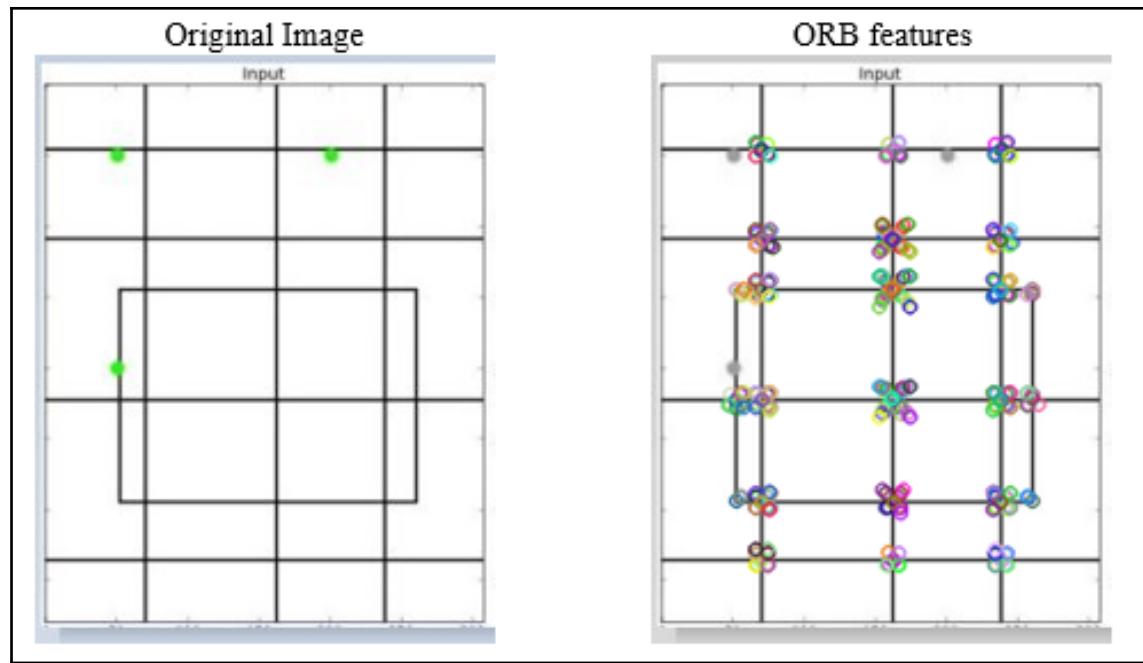




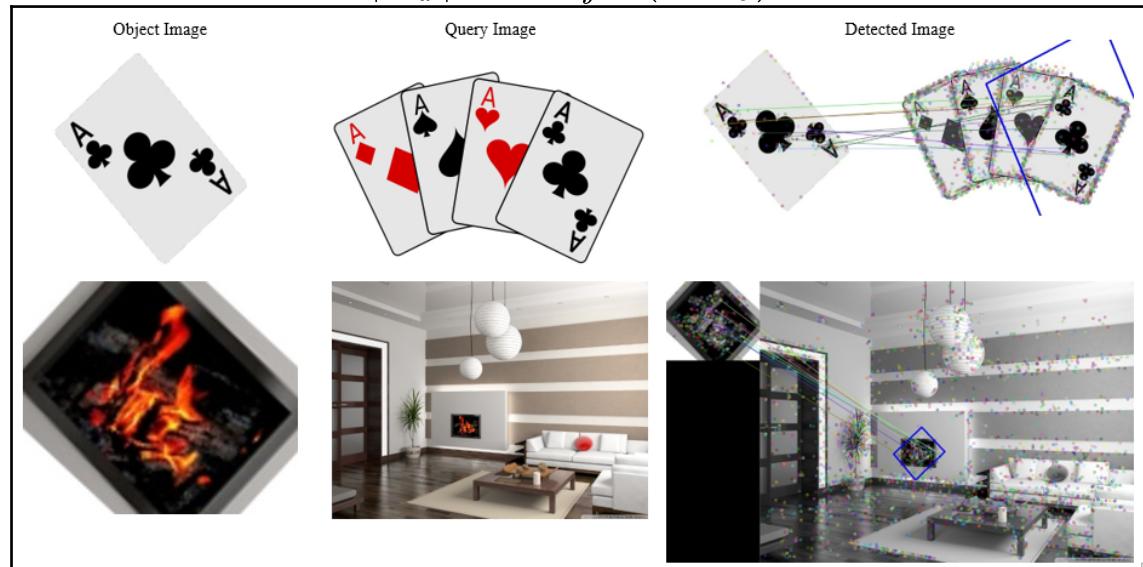
```
bhaumik@bhaumik-Lenovo-ideapad-520-15IKB:~/Desktop/opencv/chapter 7$ ./hough_line
CPU Time : 4.01799 ms
CPU FPS : 248.88
GPU Time : 1.5828 ms
GPU FPS : 631.791
```



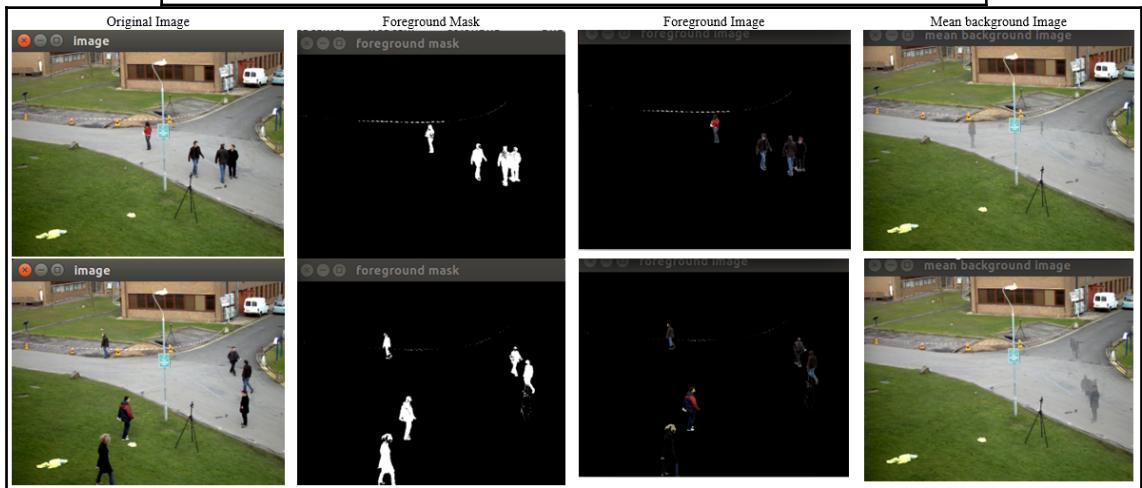


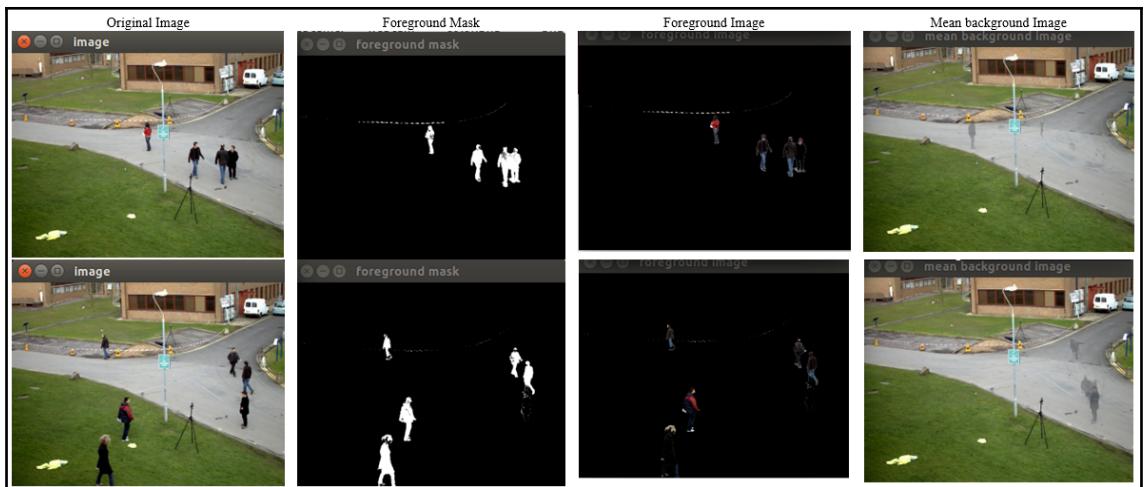


$$|H_a| = D_x D_y - (wDxy)^2$$

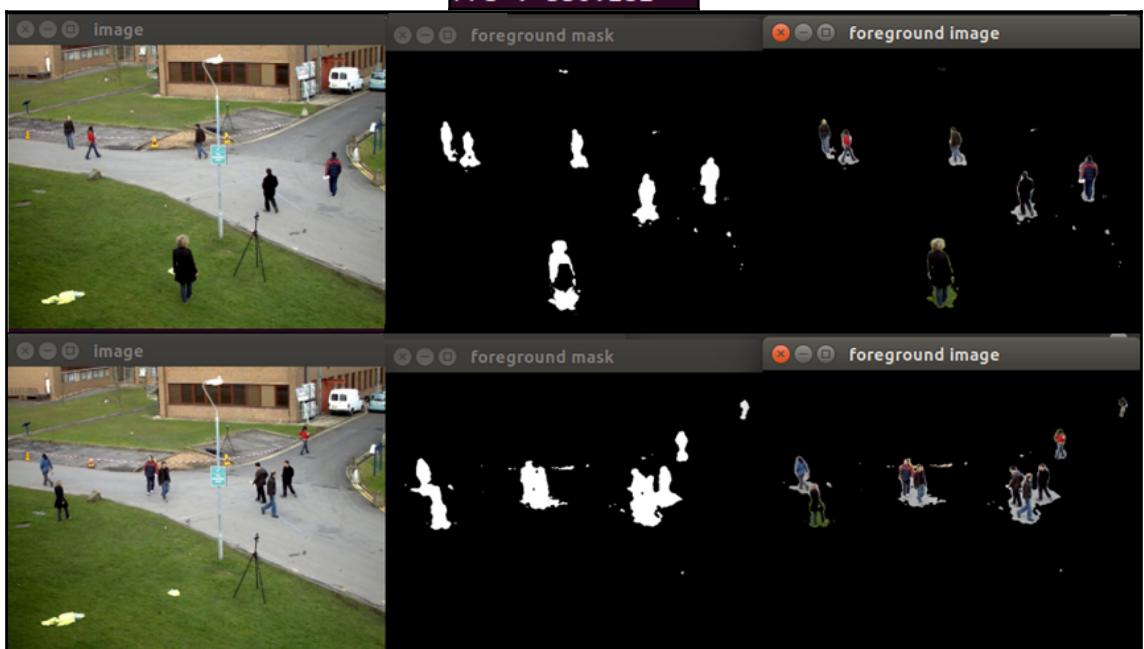








FPS : 333.944  
FPS : 332.342  
FPS : 331.221  
FPS : 330.282

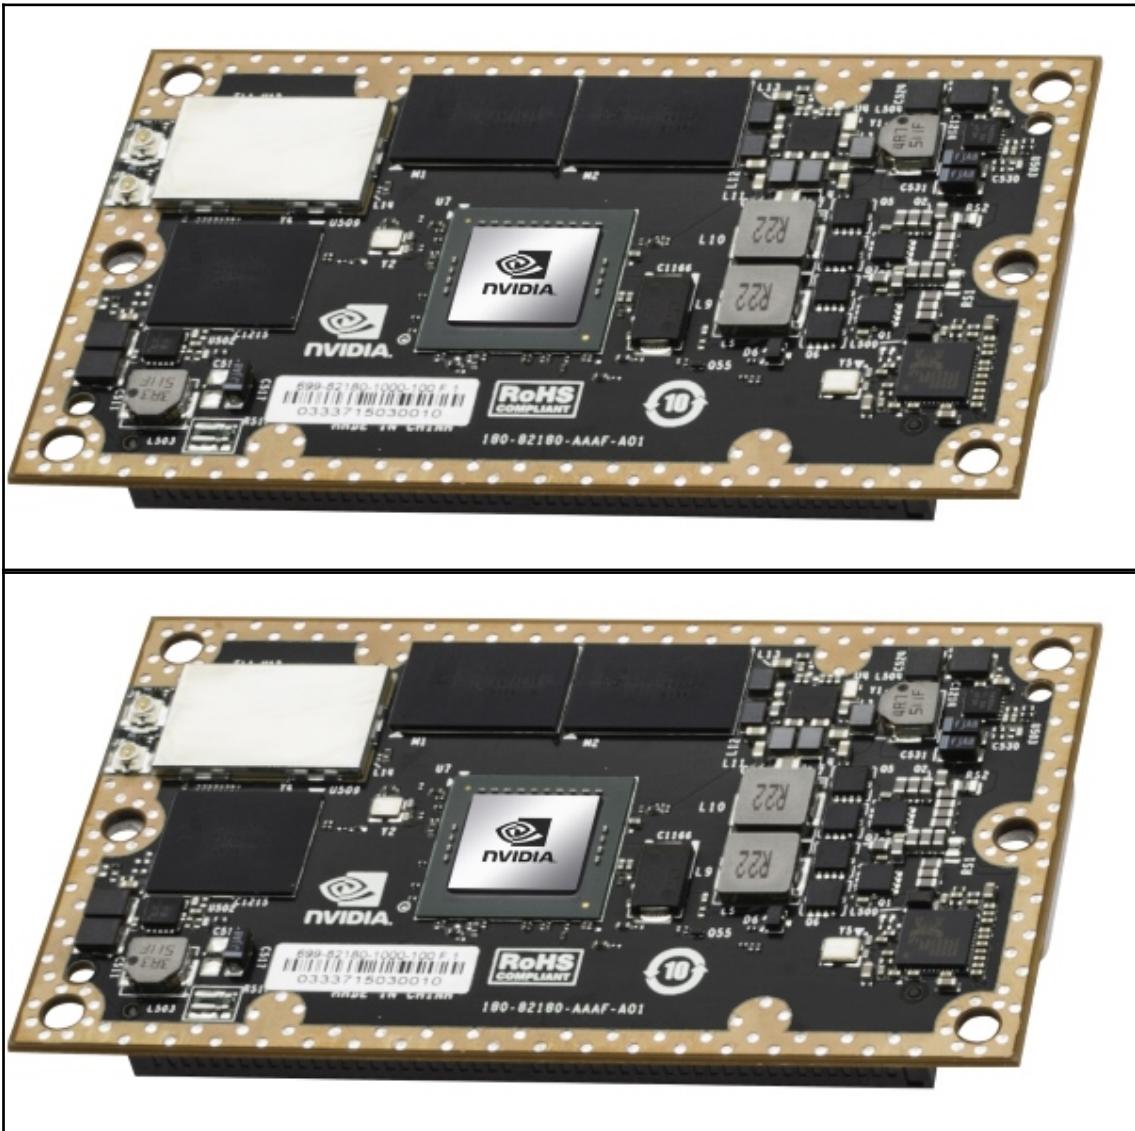


---

FPS : 112.059
FPS : 113.165
FPS : 111.566
FPS : 113.295
FPS : 113.228

---

# Chapter 8: Introduction to the Jetson TX1 Development Board and Installing OpenCV on Jetson TX1





## JetPack

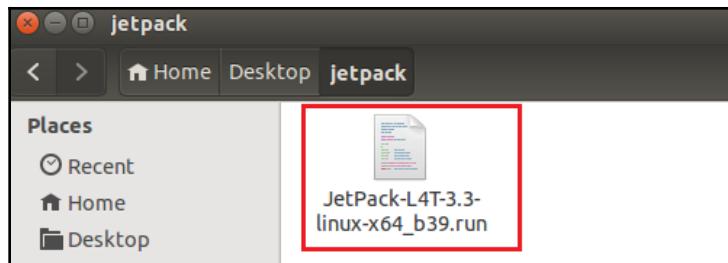
NVIDIA JetPack SDK is the most comprehensive solution for building AI applications. Use the JetPack installer to flash your Jetson Developer Kit with the latest OS image, to install developer tools for both host PC and Developer Kit, and to install the libraries and APIs, samples, and documentation needed to jumpstart your development environment.

JetPack 3.3 with the latest BSPs [[L4T 28.2.1 for Jetson TX2/TX2i](#) and [L4T 28.2 for Jetson TX1](#)] is the latest production software release for NVIDIA Jetson TX2, Jetson TX2i, and Jetson TX1. It bundles all the Jetson platform software, including TensorRT, cuDNN, CUDA Toolkit, VisionWorks, GStreamer, and OpenCV, all built on top of L4T with LTS Linux kernel.

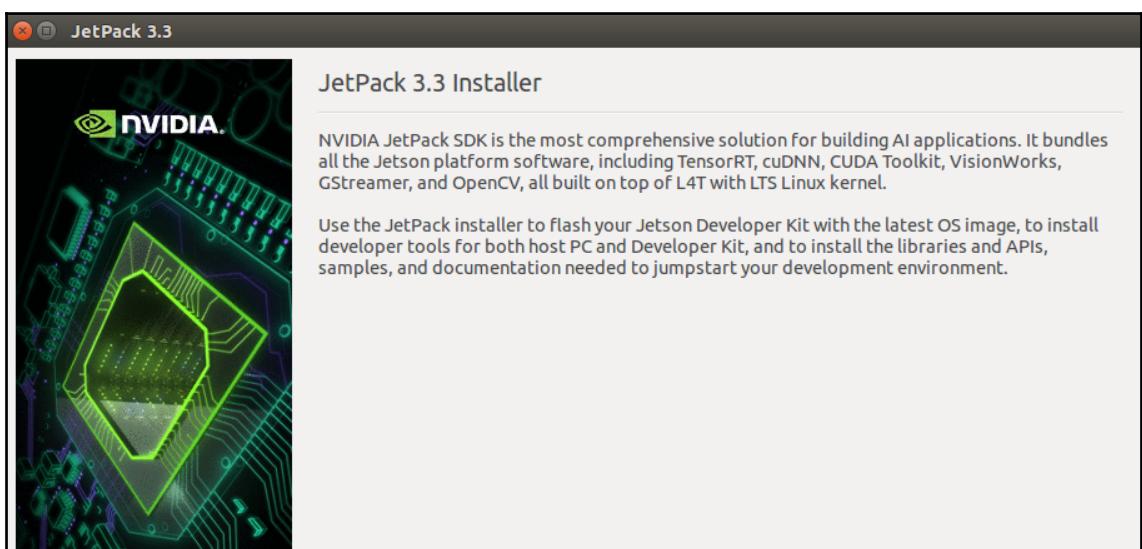
The highlight of this release is TensorRT 4.0, enabling support for TensorFlow's TensorRT integration feature. Additionally, cuDNN has a small point release to support the new TensorRT version, while all other JetPack components remain unchanged from JetPack 3.2.1.

[View the full 3.3 Release Notes here.](#)

 [Download Jetpack](#)



```
bhaumik@bhaumik-VirtualBox:~$ cd Desktop/jetpack
bhaumik@bhaumik-VirtualBox:~/Desktop/jetpack$ chmod +x JetPack-L4T-3.3-linux-x64
_b39.run
bhaumik@bhaumik-VirtualBox:~/Desktop/jetpack$ ./JetPack-L4T-3.3-linux-x64_b39.ru
n
Creating directory _installer
Verifying archive integrity... All good.
Uncompressing JetPack    57% ■
```



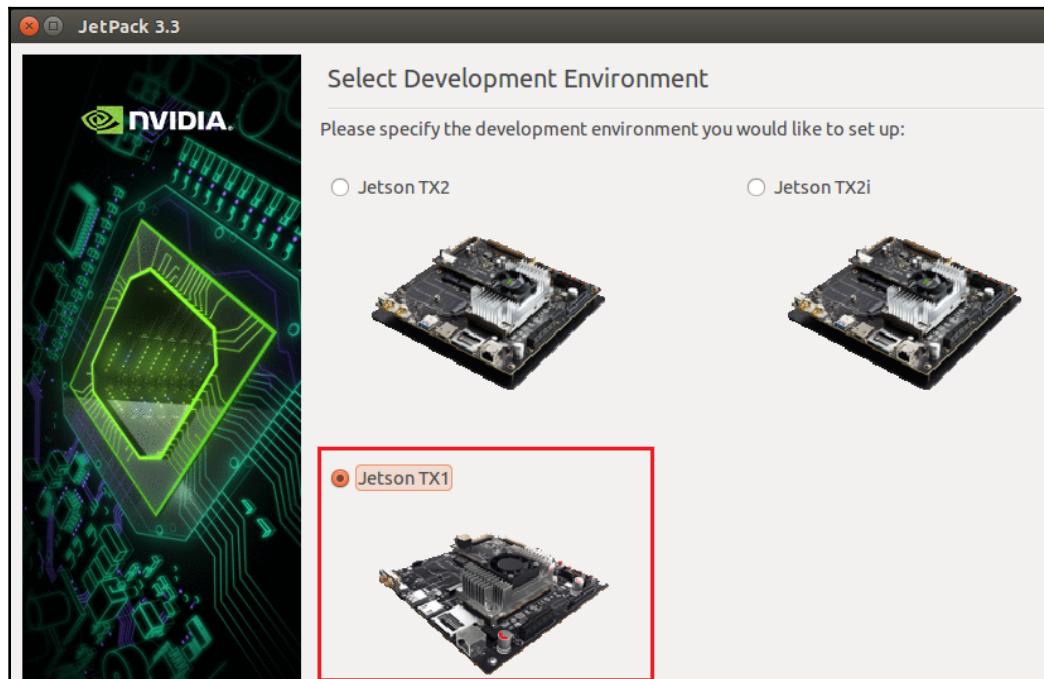
## Installation Configuration

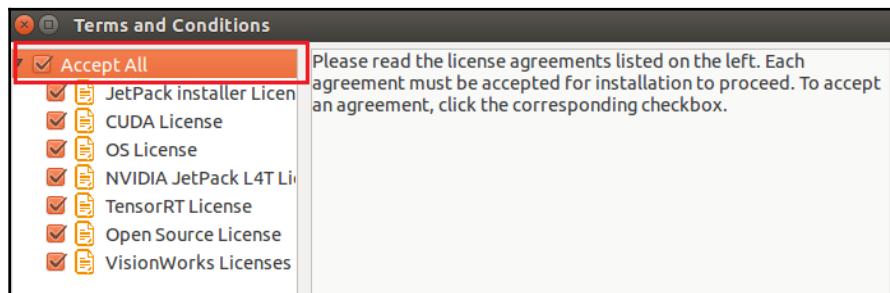
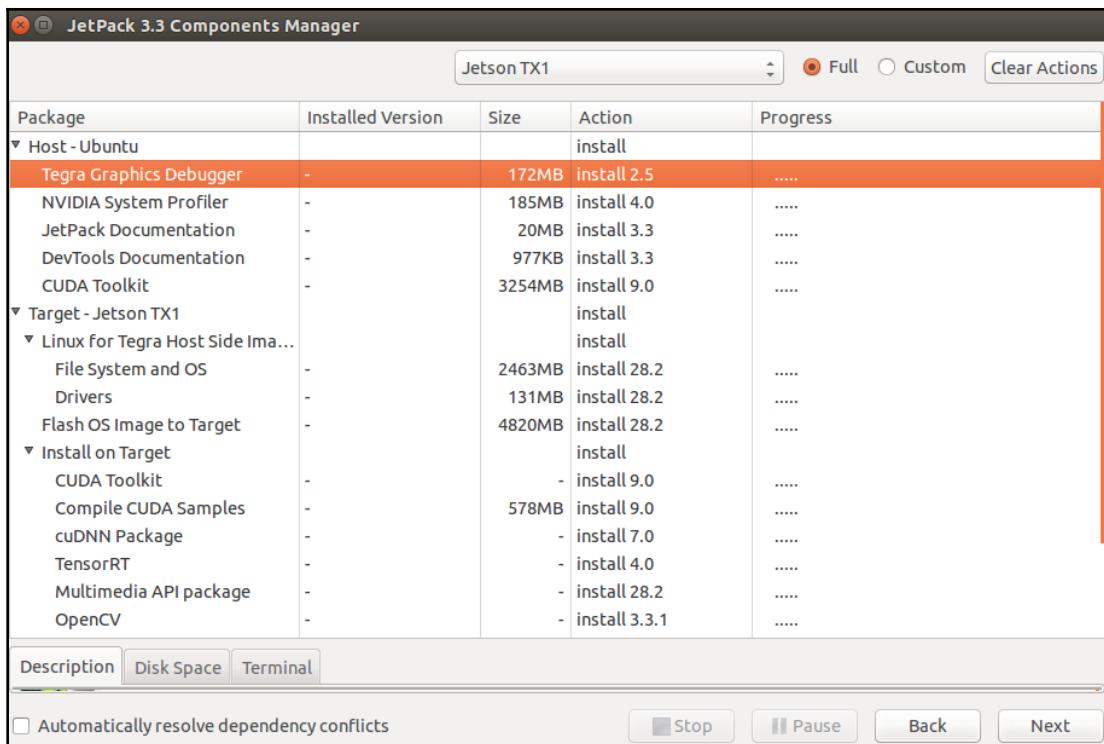
Please specify the directory where JetPack 3.3 will be installed.

Installation Directory:  ...

Please specify the directory where the components will be downloaded.

Download Directory:  ...





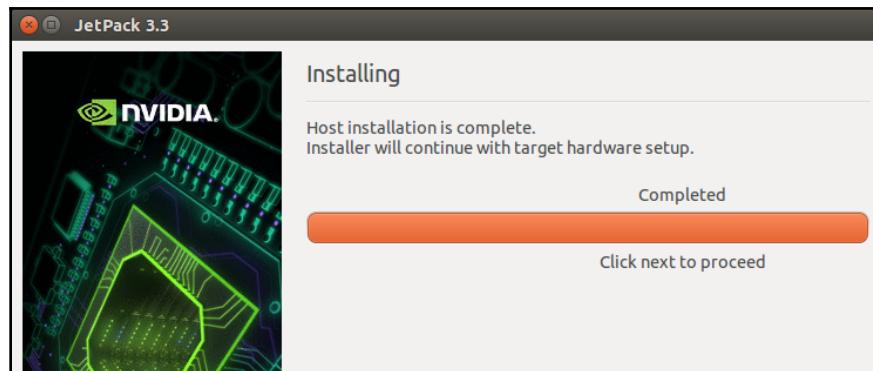
JetPack 3.3 Components Manager

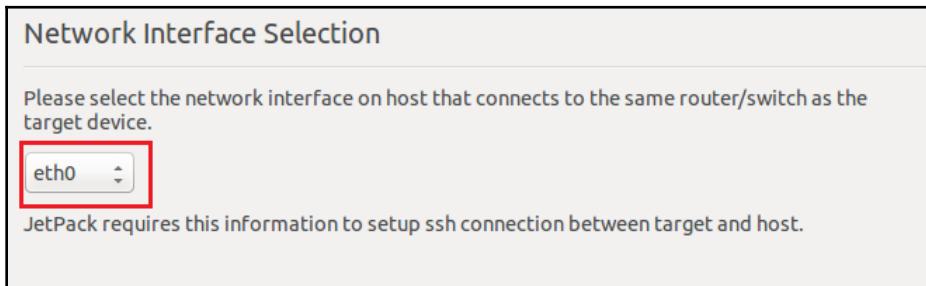
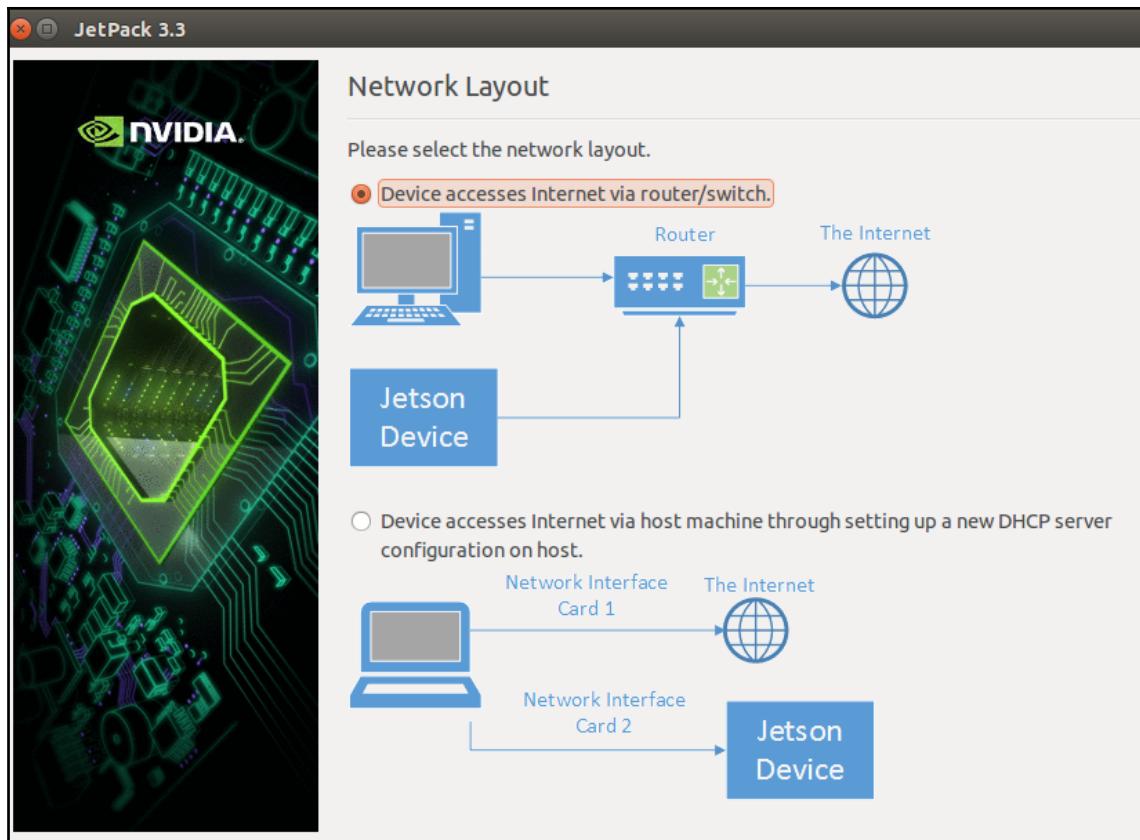
Jetson TX1 Full Custom Clear Actions

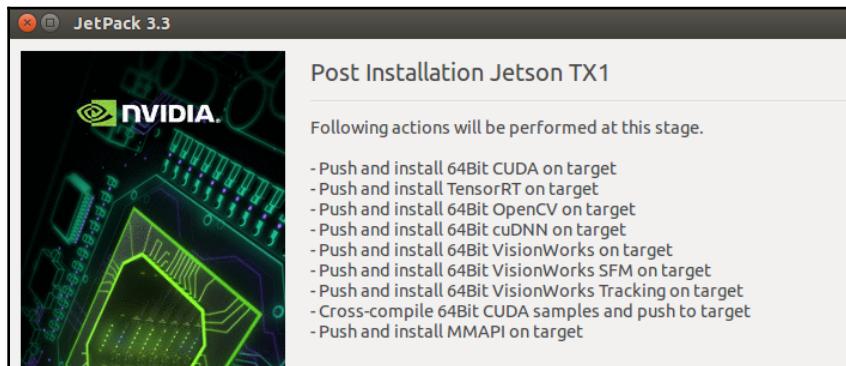
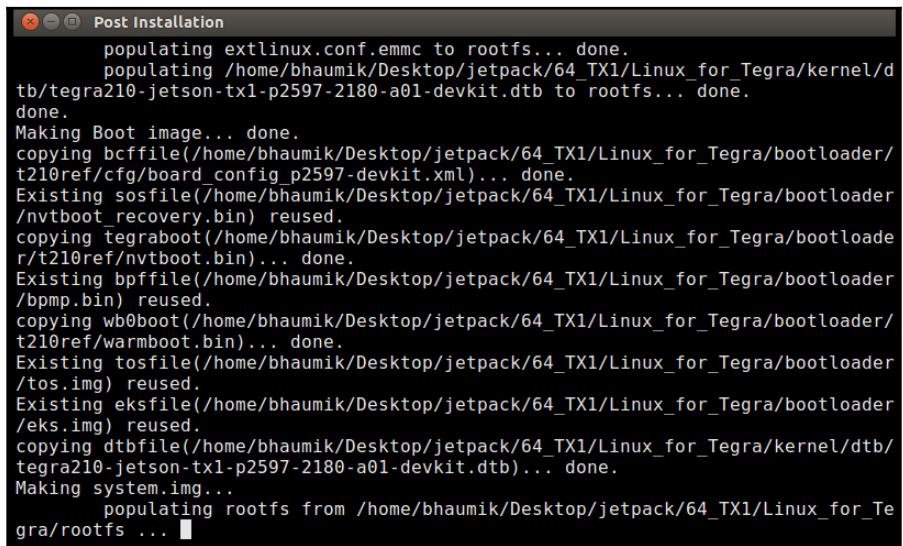
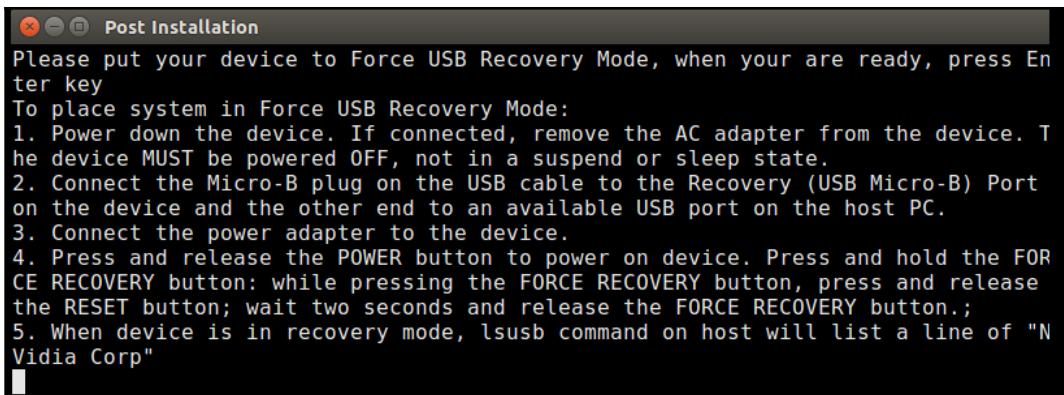
Package	Installed Version	Size	Action	Progress
Host - Ubuntu			install	
Tegra Graphics Debugger	-	172MB	install 2.5	5m 50s remaining - 6.5% (754 KB/s)
NVIDIA System Profiler	-	185MB	install 4.0	3m 7s remaining - 7.2% (441 KB/s)
JetPack Documentation	-	20MB	install 3.3	1m 20s remaining - 66.7% (7 KB/s)
DevTools Documentation	-	977KB	install 3.3	....
CUDA Toolkit	-	3254MB	install 9.0	....
Target - Jetson TX1			install	
Linux for Tegra Host Side Image			install	
File System and OS	-	2463MB	install 28.2	....
Drivers	-	131MB	install 28.2	....
Flash OS Image to Target	-	4820MB	install 28.2	....
Install on Target			install	
CUDA Toolkit	-	-	install 9.0	....
Compile CUDA Samples	-	578MB	install 9.0	....
cuDNN Package	-	-	install 7.0	....
TensorRT	-	-	install 4.0	....
Multimedia API package	-	-	install 28.2	....
OpenCV	-	-	install 3.3.1	....

Description Disk Space Terminal

Automatically resolve dependency conflicts Stop Pause Back Next

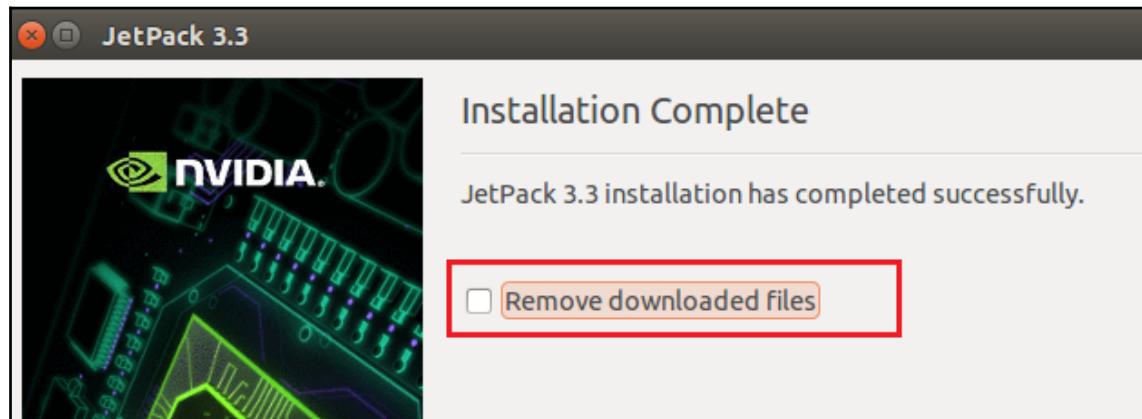






---

```
Installing CUDA on target
Copying /home/bhaumik/Desktop/jetpack/jetpack_download/cuda-repo-l4t-9-0-local_9
.0.252-1_arm64.deb file to target...
cuda-repo-l4t-9-0-local_9.0.252-1_arm64.deb
 343,179,264  56%   7.07MB/s   0:00:36
```



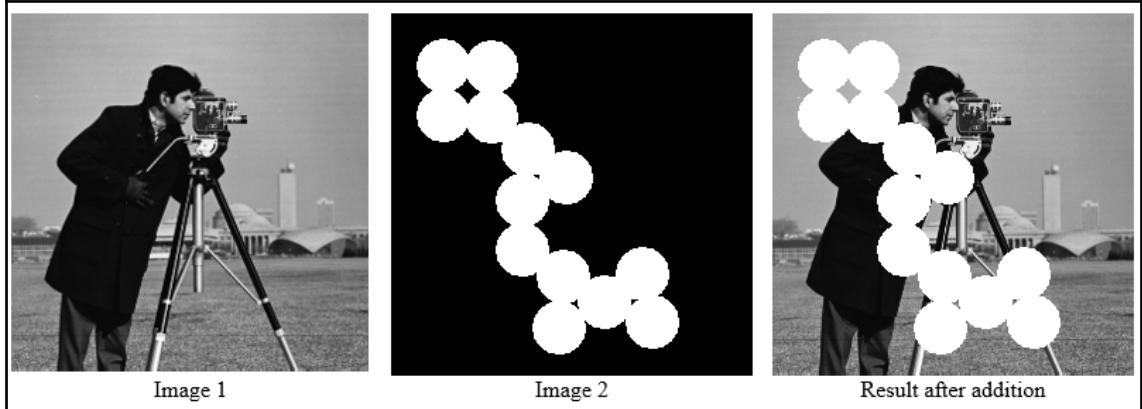
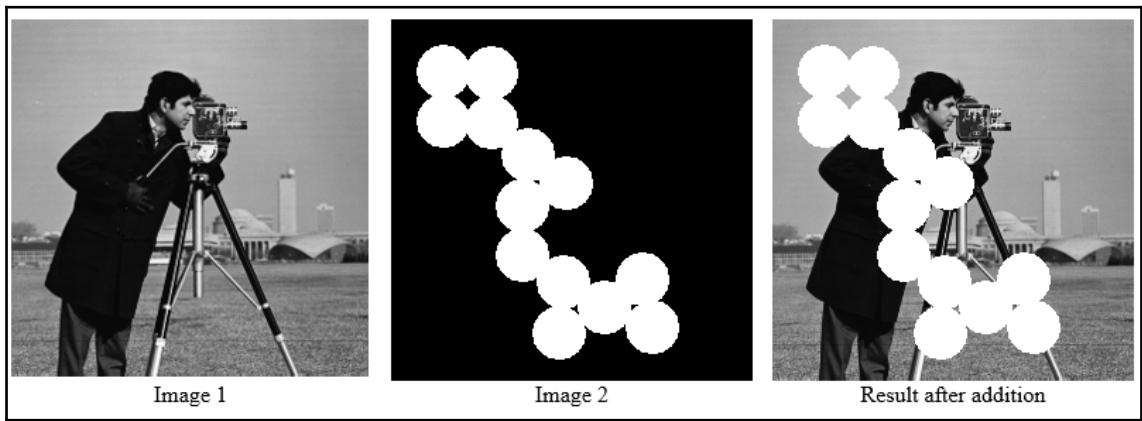
# Chapter 9: Deploying Computer Vision Applications on Jetson TX1

```
ubuntu@tegra-ubuntu:~/Desktop/opencv/Chapter 9$ nvcc kernel.cu -o device
ubuntu@tegra-ubuntu:~/Desktop/opencv/Chapter 9$ ./device
./device Starting...
CUDA Device Query (Runtime API) version (CUDART static linking)

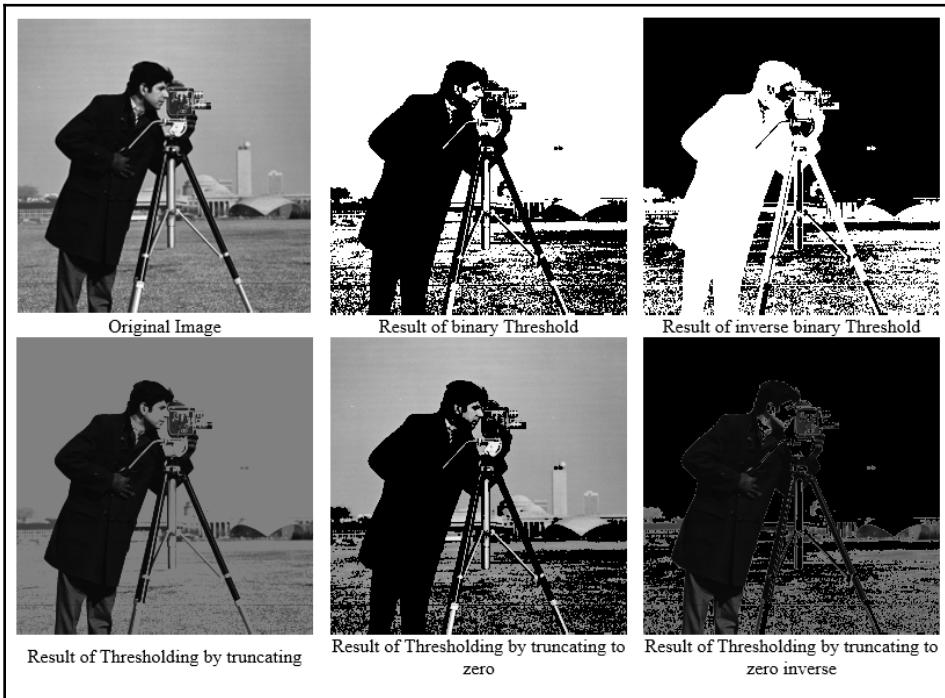
Detected 1 CUDA Capable device(s)

Device 0: "NVIDIA Tegra X1"
  CUDA Driver Version / Runtime Version      9.0 / 9.0
  CUDA Capability Major/Minor version number: 5.3
  Total amount of global memory:            3984 MBytes (4177342464 bytes)
    ( 2) Multiprocessors GPU Max Clock rate:          998 MHz (1.00 GHz)
  Memory Clock rate:                      13 Mhz
  Memory Bus Width:                      64-bit
  L2 Cache Size:                         262144 bytes
  Maximum Texture Dimension Size (x,y,z)   1D=(65536), 2D=(65536, 65536), 3D=(4096, 4096, 4096)
  Maximum Layered 1D Texture Size, (num) layers  1D=(16384), 2048 layers
  Maximum Layered 2D Texture Size, (num) layers  2D=(16384, 16384), 2048 layers
  Total amount of constant memory:          65536 bytes
  Total amount of shared memory per block:  49152 bytes
  Total number of registers available per block: 32768
  Warp size:                            32
  Maximum number of threads per multiprocessor: 2048
  Maximum number of threads per block:       1024
  Max dimension size of a thread block (x,y,z): (1024, 1024, 64)
  Max dimension size of a grid size (x,y,z): (2147483647, 65535, 65535)
  Maximum memory pitch:                  2147483647 bytes
  Texture alignment:                   512 bytes
  Concurrent copy and kernel execution: Yes with 1 copy engine(s)
  Run time limit on kernels:           Yes
  Integrated GPU sharing Host Memory: Yes
  Support host page-locked memory mapping: Yes
  Alignment requirement for Surfaces: Yes
  Device has ECC support:             Disabled
  Device supports Unified Addressing (UVA): Yes
  Supports Cooperative Kernel Launch: No
  Supports MultiDevice Co-op Kernel Launch: No
  Device PCI Domain ID / Bus ID / location ID: 0 / 0 / 0
  Compute Mode:
    < Default (multiple host threads can use ::cudaSetDevice() with device simultaneously) >
ubuntu@tegra-ubuntu:~/Desktop/opencv/Chapter 9$
```

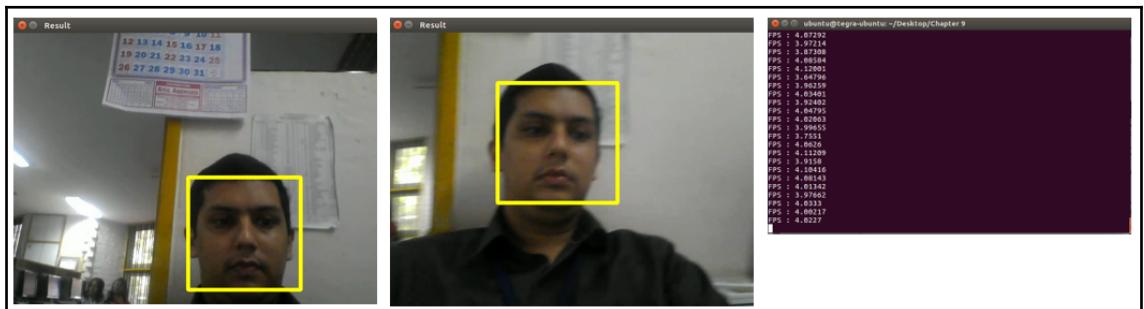
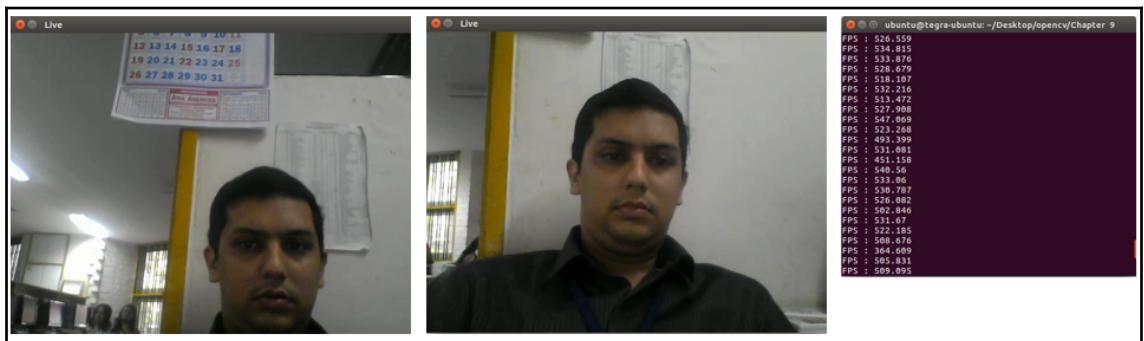
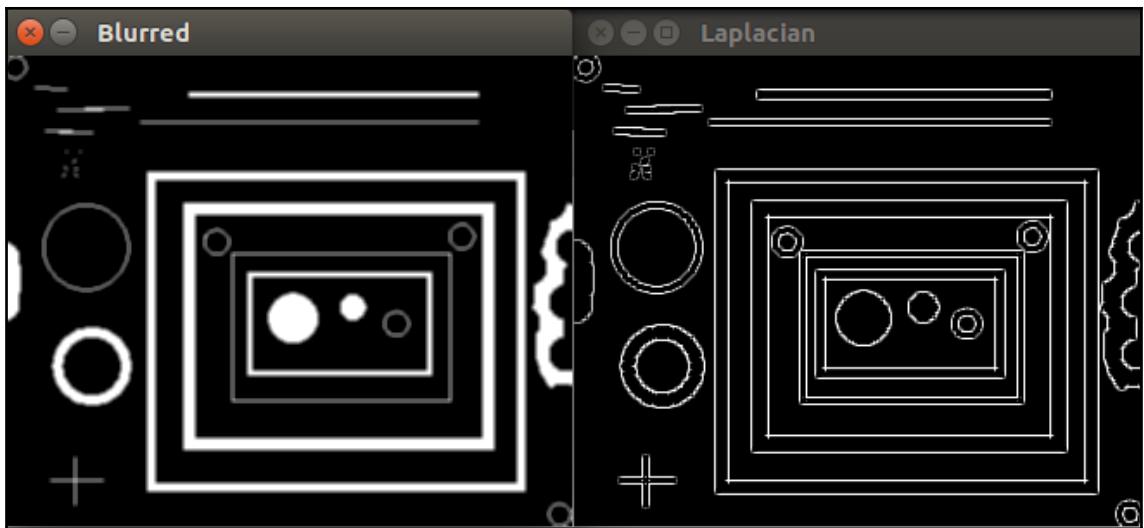
```
ubuntu@tegra-ubuntu:~/Desktop/opencv/Chapter 9$ nvcc 01_performance_cuda_events
.cu -o gpu_add
ubuntu@tegra-ubuntu:~/Desktop/opencv/Chapter 9$ ./gpu_add
Time to add 50000 numbers: 3.4 ms
Vector addition on GPU
GPU has computed Sum Correctly
```

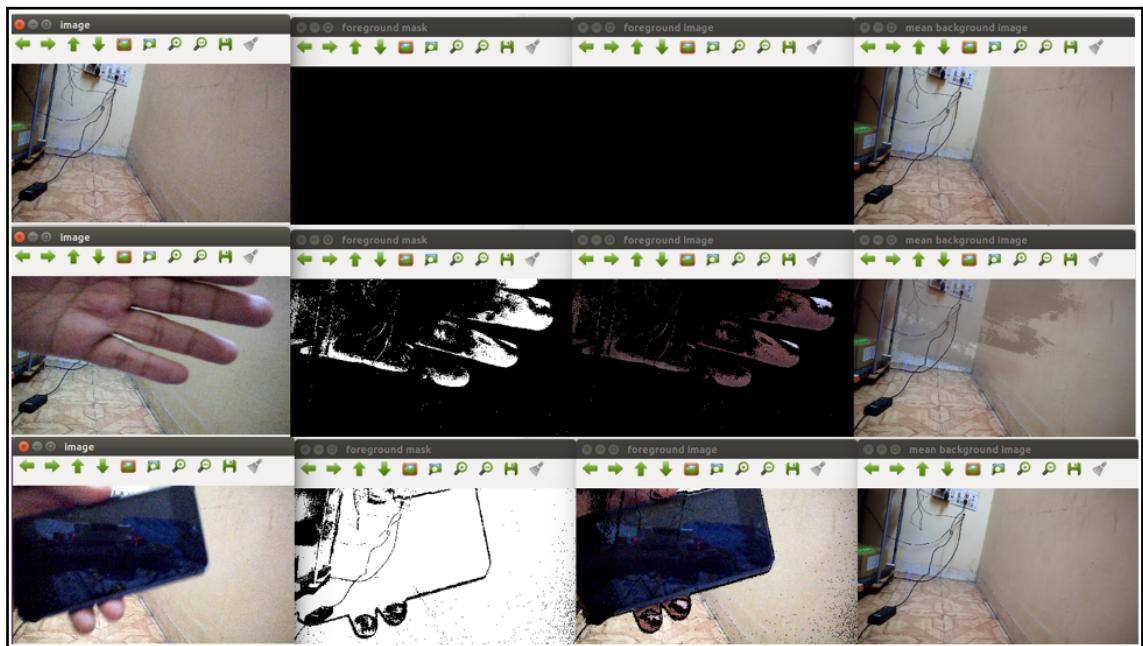


```
ubuntu@tegra-ubuntu:~/Desktop/opencv/Chapter 9$ ./image_add
Performance of Addition on Jetson TX!:
Time: 0.000262104
FPS: 3815.28
```



```
Performance of Thresholding
Time: 0.000326671
FPS: 3061.18
```





---

```
ubuntu@tegra-ubuntu: ~/Desktop/Chapter 9
FPS : 45.9451
FPS : 38.3942
FPS : 39.1694
FPS : 39.365
FPS : 38.7298
FPS : 34.2839
FPS : 75.2029
FPS : 59.0898
FPS : 47.6766
FPS : 49.3542
FPS : 34.7506
FPS : 38.3625
FPS : 38.4274
FPS : 44.3814
FPS : 77.2785
FPS : 65.9443
FPS : 58.8832
FPS : 54.1984
FPS : 49.1819
FPS : 78.8699
FPS : 71.9925
FPS : 65.1592
FPS : 60.9173
```



# Chapter 10: Getting Started with PyCUDA

**Select Target Platform **

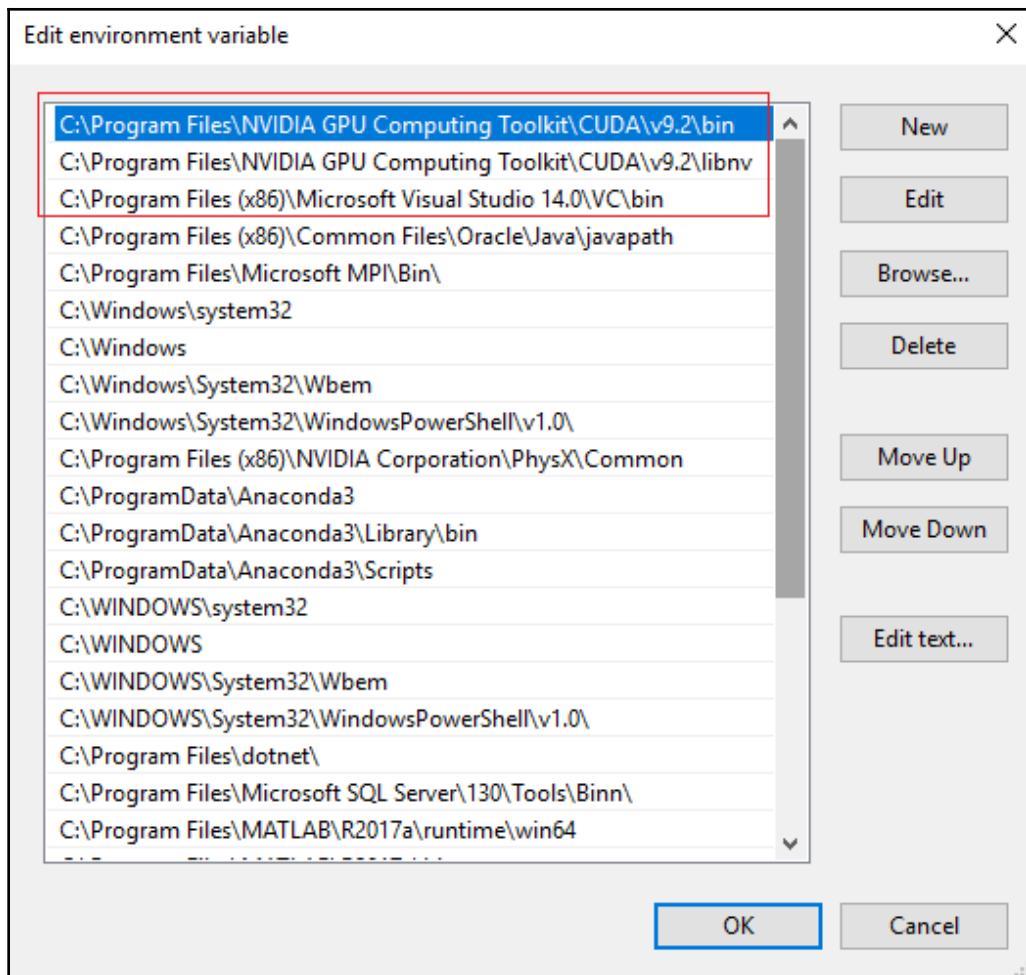
Click on the green buttons that describe your target platform. Only supported platforms will be shown.

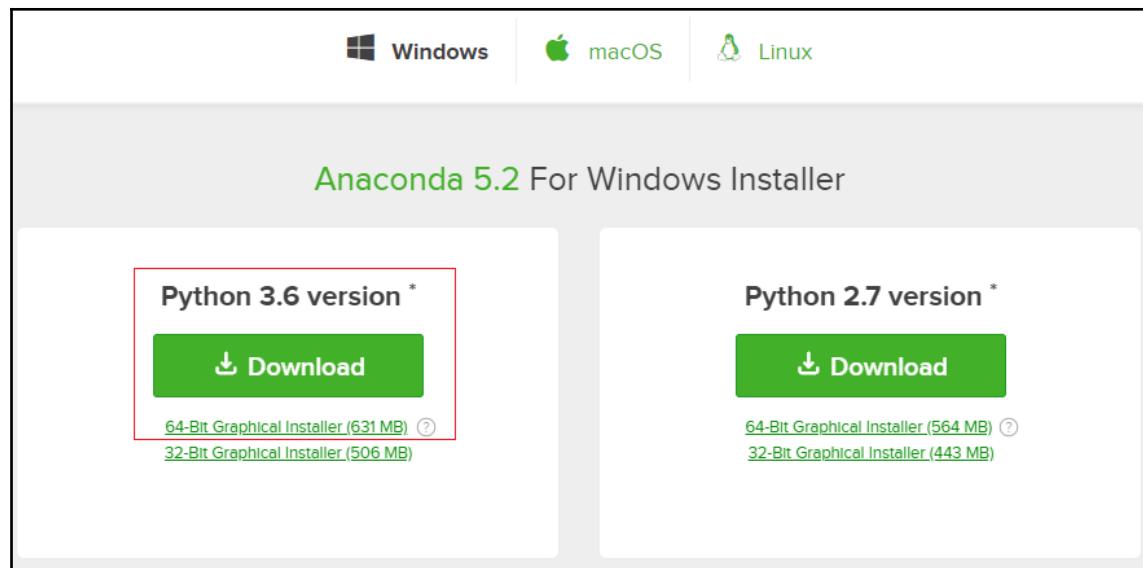
Operating System	<input type="button" value="Windows"/> <input type="button" value="Linux"/> <input type="button" value="Mac OS X"/>
Architecture 	<input type="button" value="x86_64"/>
Version	<input type="button" value="10"/> <input type="button" value="8.1"/> <input type="button" value="7"/> <input type="button" value="Server 2016"/> <input type="button" value="Server 2012 R2"/>
Installer Type 	<input type="button" value="exe (network)"/> <input type="button" value="exe (local)"/>

**Download Installer for Windows 10 x86\_64**

The base installer is available for download below.

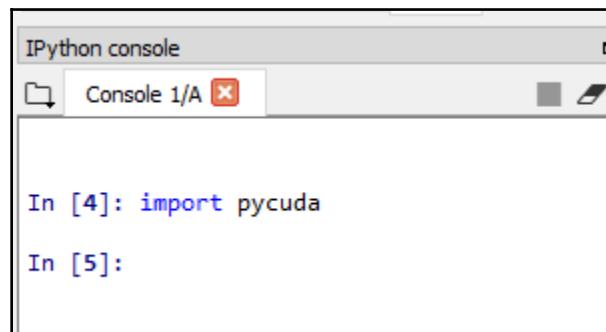
 <b>Base Installer</b>	<a href="#">Download (1.5 GB) </a>
Installation Instructions:	
<ol style="list-style-type: none"><li>1. Double click cuda_9.2.148_win10.exe</li><li>2. Follow on-screen prompts</li></ol>	





[pycuda-2017.1.1+cuda9185-cp34-cp34m-win32.whl](#)  
[pycuda-2017.1.1+cuda9185-cp34-cp34m-win\\_amd64.whl](#)  
[pycuda-2017.1.1+cuda9185-cp35-cp35m-win32.whl](#)  
[pycuda-2017.1.1+cuda9185-cp35-cp35m-win\\_amd64.whl](#)  
[pycuda-2017.1.1+cuda9185-cp36-cp36m-win32.whl](#)  
[pycuda-2017.1.1+cuda9185-cp36-cp36m-win\\_amd64.whl](#)  
[pycuda-2018.1+cuda92148-cp27-cp27m-win32.whl](#)  
[pycuda-2018.1+cuda92148-cp27-cp27m-win\\_amd64.whl](#)  
[pycuda-2018.1+cuda92148-cp34-cp34m-win32.whl](#)  
[pycuda-2018.1+cuda92148-cp34-cp34m-win\\_amd64.whl](#)  
[pycuda-2018.1+cuda92148-cp35-cp35m-win32.whl](#)  
[pycuda-2018.1+cuda92148-cp35-cp35m-win\\_amd64.whl](#)  
[pycuda-2018.1+cuda92148-cp36-cp36m-win32.whl](#)  
[pycuda-2018.1+cuda92148-cp36-cp36m-win\\_amd64.whl](#)  
[pycuda-2018.1+cuda92148-cp37-cp37m-win32.whl](#)  
[pycuda-2018.1+cuda92148-cp37-cp37m-win\\_amd64.whl](#)

```
PS C:\Users\bhaum\Downloads> pip install .\pycuda-2017.1.1+cuda92148-cp36-cp36m-win_amd64.whl
Processing c:\users\bhaum\downloads\pycuda-2017.1.1+cuda92148-cp36-cp36m-win_amd64.whl
Requirement already satisfied: pytest>=2 in c:\programdata\anaconda3\lib\site-packages (from pycuda==2017.1.1+cuda92148)
Requirement already satisfied: pytools>=2011.2 in c:\programdata\anaconda3\lib\site-packages (from pycuda==2017.1.1+cuda92148)
Requirement already satisfied: decorator>=3.2.0 in c:\programdata\anaconda3\lib\site-packages (from pycuda==2017.1.1+cuda92148)
Requirement already satisfied: appdirs>=1.4.0 in c:\programdata\anaconda3\lib\site-packages (from pycuda==2017.1.1+cuda92148)
Requirement already satisfied: py>=1.5.0 in c:\programdata\anaconda3\lib\site-packages (from pytest>=2->pycuda==2017.1.1+cuda92148)
Requirement already satisfied: six>=1.10.0 in c:\programdata\anaconda3\lib\site-packages (from pytest>=2->pycuda==2017.1.1+cuda92148)
Requirement already satisfied: setuptools in c:\programdata\anaconda3\lib\site-packages (from pytest>=2->pycuda==2017.1.1+cuda92148)
Requirement already satisfied: attrs>=17.2.0 in c:\programdata\anaconda3\lib\site-packages (from pytest>=2->pycuda==2017.1.1+cuda92148)
Requirement already satisfied: pluggy<0.7,>=0.5 in c:\programdata\anaconda3\lib\site-packages (from pytest>=2->pycuda==2017.1.1+cuda92148)
Requirement already satisfied: colorama in c:\programdata\anaconda3\lib\site-packages (from pytest>=2->pycuda==2017.1.1+cuda92148)
Requirement already satisfied: numpy>=1.6.0 in c:\programdata\anaconda3\lib\site-packages (from pytools>=2011.2->pycuda==2017.1.1+cuda92148)
Installing collected packages: pycuda
  Found existing installation: pycuda 2017.1.1+cuda9185
    Uninstalling pycuda-2017.1.1+cuda9185:
      Successfully uninstalled pycuda-2017.1.1+cuda9185
Successfully installed pycuda-2017.1.1+cuda92148
You are using pip version 9.0.1, however version 18.0 is available.
You should consider upgrading via the 'python -m pip install --upgrade pip' command.
```



## Select Target Platform

Click on the green buttons that describe your target platform. Only supported platforms will be shown.

Operating System

Windows  Linux  Mac OSX

Architecture 

x86\_64  ppc64le

Distribution

Fedora  OpenSUSE  RHEL  CentOS  SLES  Ubuntu

Version

17.10  16.04

Installer Type 

runfile (local)  deb (local)  deb (network)

## Download Installer for Linux Ubuntu 17.10 x86\_64

The base installer is available for download below.

 Base Installer

[Download \[1.7 GB\] !\[\]\(71e0e0275eeca8815cc295413b2e7f4b\_img.jpg\)](#)

Installation Instructions:

1. Run `sudo sh cuda\_9.2.148\_396.37\_linux.run`
2. Follow the command-line prompts



Windows



macOS



Linux

## Anaconda 5.2 For Linux Installer

### Python 3.6 version \*

Download

[64-Bit \(x86\) Installer \(622 MB\)](#) (?)  
[64-Bit \(Power8 and Power9\) Installer \(288 MB\)](#)  
[32-Bit Installer \(507 MB\)](#)

### Python 2.7 version \*

Download

[64-Bit \(x86\) Installer \(603 MB\)](#) (?)  
[64-Bit \(Power8 and Power9\) Installer \(270 MB\)](#)  
[32-Bit Installer \(489 MB\)](#)

```
bhaumik@bhaumik-Lenovo-ideapad-520-15IKB:~$ conda install -c lukepfister pycuda
Solving environment: done
```

```
==> WARNING: A newer version of conda exists. <==
  current version: 4.4.10
  latest version: 4.5.8
```

```
Please update conda by running
```

```
$ conda update -n base conda
```

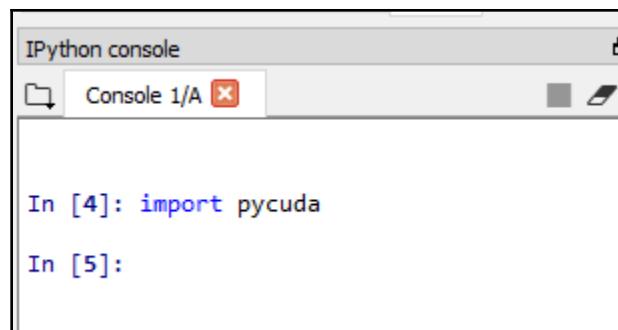
```
## Package Plan ##
```

```
environment location: /home/bhaumik/anaconda3
```

```
added / updated specs:
  - pycuda
```

```
The following packages will be downloaded:
```

package	build		
ca-certificates-2018.03.07	0	124 KB	
pycuda-2017.1	py36_0	627 KB	lukepfister
appdirs-1.4.3	py36h28b3542_0	16 KB	
pytools-2018.4	py36_0	110 KB	lukepfister
mako-1.0.4	py36_0	116 KB	lukepfister
openssl-1.0.2o	h20670df_0	3.4 MB	
certifi-2018.4.16	py36_0	142 KB	
Total:		4.6 MB	



# Chapter 11: Working with PyCUDA

```
PS G:\cuda opencv book material\CUDA book code\Chapter11> python .\hello_pycuda.py
Hello,PyCUDA!!!
PS G:\cuda opencv book material\CUDA book code\Chapter11> _
```

A screenshot of a Windows PowerShell window titled "Windows PowerShell". The window displays detailed GPU configuration information for a GeForce 940MX. The output includes various parameters such as Device ID, Compute Capability, Total Memory, Async Engine Count, Clock Rate, Compute Capability Minor, Compute Mode, Concurrent Kernels, ECC Enabled, Global Memory Cache Supported, Global Memory Bus Width, GPU Overlap, Integrated, Kernel Exec Timeout, Local L1 Cache Supported, Managed Memory, Maximum Surface 1D Layered Layers, Maximum Surface 1D Width, Maximum Surface 2D Layered Height, Maximum Surface 2D Layered Layers, Maximum Surface 2D Layered Width, Maximum Surface 3D Depth, Maximum Surface 3D Width, Maximum Surface 3D Height, Maximum Surface Cubemap Layered Layers, Maximum Surface Cubemap Width, Maximum Surface Cubemap Layered Width, Maximum Texture 1D Layered Layers, Maximum Texture 1D Linear Width, Maximum Texture 1D Mipmapped Width, Maximum Texture 1D Width, Maximum Texture 2D Array Height, Maximum Texture 2D Array Width, Maximum Texture 2D Gather Height, Maximum Texture 2D Gather Width, Maximum Texture 2D Linear Height, Maximum Texture 2D Linear Width, Maximum Texture 2D Mipmapped Width, Maximum Texture 2D Width, Maximum Texture 3D Depth, Maximum Texture 3D Height, Maximum Texture 3D Width, Maximum Texture 3D Alternate, Maximum Texture 3D Height Alternate, Maximum Texture Cubemap Layered Layers, Maximum Texture Cubemap Width, Maximum Texture Cubemap Layered Width, Maximum Grid Dim X, Maximum Grid Dim Y, Maximum Grid Dim Z, Maximum Registers Per Multiprocessor, Maximum Shared Memory Per Block, Maximum Shared Memory Per Multiprocessor, Maximum Threads Per Multiprocessor, Maximum Threads Per Multiprocessor Count, Multi GPU Board, PCI Device ID, PCI Domain ID, Stream Priorities Supported, TCC Driver, and Texture Pitch Alignment. The timestamp at the bottom right of the window is 13:27 12-08-2018.

```
Device #: GeForce 940MX
Compute Capability: 5.0
Total Memory: 4 GB
ASYNC_ENGINE_COUNT: 1   CAN_MAP_HOST_MEMORY: 1
CLOCK_RATE: 1189000    COMPUTE_CAPABILITY_MAJOR: 5
COMPUTE_CAPABILITY_MINOR: 0   COMPUTE_MODE: DEFAULT
CONCURRENT_KERNELS: 0   ECC_ENABLED: 0
GLOBAL_MEMORY_CACHE_SUPPORTED: 0   GLOBAL_MEMORY_BUS_WIDTH: 64
GPU_OVERLAP: 0   INTEGRATED: 0
KERNEL_EXEC_TIMEOUT: 1   L2_CACHE_SIZE: 1048576
LOCAL_L1_CACHE_SUPPORTED: 1   MANAGED_MEMORY: 1
MAXIMUM_SURFACE1D_LAYERED_LAYERS: 2048   MAXIMUM_SURFACE1D_LAYERED_WIDTH: 16384
MAXIMUM_SURFACE1D_WIDTH: 16384   MAXIMUM_SURFACE2D_HEIGHT: 65536
MAXIMUM_SURFACE2D_LAYERED_HEIGHT: 16384   MAXIMUM_SURFACE2D_LAYERED_LAYERS: 2048
MAXIMUM_SURFACE2D_LAYERED_WIDTH: 16384   MAXIMUM_SURFACE2D_WIDTH: 65536
MAXIMUM_SURFACE3D_DEPTH: 4096   MAXIMUM_SURFACE3D_HEIGHT: 4096
MAXIMUM_SURFACE3D_WIDTH: 4096   MAXIMUM_SURFACECUBEMAP_LAYERED_LAYERS: 2046
MAXIMUM_SURFACECUBEMAP_LAYERED_WIDTH: 16384   MAXIMUM_SURFACECUBEMAP_WIDTH: 16384
MAXIMUM_TEXTURE1D_LAYERED_LAYERS: 2048   MAXIMUM_TEXTURE1D_LAYERED_WIDTH: 16384
MAXIMUM_TEXTURE1D_LINEAR_WIDTH: 134217728   MAXIMUM_TEXTURE1D_MIPMAPPED_WIDTH: 16384
MAXIMUM_TEXTURE1D_WIDTH: 65536   MAXIMUM_TEXTURE2D_ARRAY_HEIGHT: 16384
MAXIMUM_TEXTURE2D_ARRAY_NUMSLICES: 2048   MAXIMUM_TEXTURE2D_ARRAY_WIDTH: 16384
MAXIMUM_TEXTURE2D_GATHER_HEIGHT: 16384   MAXIMUM_TEXTURE2D_GATHER_WIDTH: 16384
MAXIMUM_TEXTURE2D_LINEAR_HEIGHT: 65536   MAXIMUM_TEXTURE2D_LINEAR_WIDTH: 65536
MAXIMUM_TEXTURE2D_MIPMAPPED_HEIGHT: 16384   MAXIMUM_TEXTURE2D_MIPMAPPED_WIDTH: 16384
MAXIMUM_TEXTURE2D_WIDTH: 65536   MAXIMUM_TEXTURE3D_DEPTH: 4096
MAXIMUM_TEXTURE3D_DEPTH_ALTERNATE: 16384   MAXIMUM_TEXTURE3D_HEIGHT: 4096
MAXIMUM_TEXTURE3D_HEIGHT_ALTERNATE: 2048   MAXIMUM_TEXTURE3D_WIDTH: 4096
MAXIMUM_TEXTURE3D_WIDTH_ALTERNATE: 2048   MAXIMUM_TEXTURECUBEMAP_LAYERED_LAYERS: 2046
MAXIMUM_TEXTURECUBEMAP_LAYERED_WIDTH: 16384   MAXIMUM_TEXTURECUBEMAP_WIDTH: 16384
MAX_BLOCK_DIM_X: 1024   MAX_BLOCK_DIM_Y: 1024
MAX_BLOCK_DIM_Z: 64   MAX_GRID_DIM_X: 2147483647
MAX_GRID_DIM_Y: 65535   MAX_GRID_DIM_Z: 65535
MAX_PITCH: 2147483647   MAX_REGISTERS_PER_BLOCK: 65536
MAX_REGISTERS_PER_MULTIPROCESSOR: 65536   MAX_SHARED_MEMORY_PER_BLOCK: 49152
MAX_SHARED_MEMORY_PER_MULTIPROCESSOR: 65536   MAX_THREADS_PER_BLOCK: 1024
MAX_THREADS_PER_MULTIPROCESSOR: 2048   MEMORY_CLOCK_RATE: 2505000
MULTI_GPU_BOARD_COUNT: 3   MULTI_GPU_BOARD: 0
MULTI_GPU_BOARD_GROUP_ID: 0   PCI_BUS_ID: 1
PCI_DEVICE_ID: 0   PCI_DOMAIN_ID: 0
STREAM_PRIORITIES_SUPPORTED: 1   SURFACE_ALIGNMENT: 512
TCC_DRIVER: 0   TEXTURE_ALIGNMENT: 512
TEXTURE_PITCH_ALIGNMENT: 32   TOTAL_CONSTANT_MEMORY: 65536
```

A screenshot of a Windows PowerShell window titled "Windows PowerShell". The window displays the output of a Python script named "thread\_execution.py". The output shows the message "I am in block no: 2" repeated 10 times, followed by the message "PS G:\cuda opencv book material\CUDA book code\Chapter11>". The timestamp at the bottom right of the window is 13:27 12-08-2018.

```
I am in block no: 2
I am in block no: 7
I am in block no: 0
I am in block no: 8
I am in block no: 1
I am in block no: 3
I am in block no: 5
I am in block no: 4
I am in block no: 6
I am in block no: 9
PS G:\cuda opencv book material\CUDA book code\Chapter11>
```

```
In [11]: runfile('G:/cuda opencv book material/CUDA book code/Chapter11/add_n.py',
wdir='G:/cuda opencv book material/CUDA book code/Chapter11')
Addition on GPU:
1.2273334 + 1.3404454 = 2.5677788
```

```
Addition on GPU:  
0.7673203 + 0.5080069 = 1.2753272  
0.28488383 + 0.15324554 = 0.43812937  
-0.3220178 + -0.10700232 = -0.4290201  
-0.501334 + -1.7047318 = -2.206066  
0.023053076 + 0.34796545 = 0.37101853  
-0.44806996 + -2.071736 = -2.5198061  
2.9193559 + 0.7721601 = 3.691516  
-0.8016763 + -0.31726292 = -1.1189392  
0.607628 + -1.2539302 = -0.6463022  
-0.80436313 + 1.2789425 = 0.47457933
```

```
In [9]: runfile('G:/cuda opencv book material/CUDA book code/Chapter11/add_number.py',  
wdir='G:/cuda opencv book material/CUDA book code/Chapter11')  
Addition of 1000000 element of GPU  
0.009422s  
Addition of 1000000 element of CPU  
0.41515421867370605 s
```

```
Time of Squaring on GPU without inout  
0.149003s  
original array:  
[[3. 2. 1. 4. 3.]  
 [1. 1. 1. 2. 2.]  
 [2. 1. 1. 3. 2.]  
 [3. 4. 2. 3. 1.]  
 [4. 1. 4. 3. 1.]]  
Square with kernel:  
[[ 9. 4. 1. 16. 9.]  
 [ 1. 1. 1. 4. 4.]  
 [ 4. 1. 1. 9. 4.]  
 [ 9. 16. 4. 9. 1.]  
 [16. 1. 16. 9. 1.]]
```

```
Square with InOut:  
[[ 9. 4. 1. 16. 9.]  
 [ 1. 1. 1. 4. 4.]  
 [ 4. 1. 1. 9. 4.]  
 [ 9. 16. 4. 9. 1.]  
 [16. 1. 16. 9. 1.]]  
Time of Squaring on GPU with inout  
0.004260s
```

```

original array:
[[1 1 3 1 4]
 [4 3 4 2 1]
 [4 1 1 4 4]
 [2 3 2 4 1]
 [2 1 2 4 3]]
Squared with gpuarray:
[[ 1.  1.  9.  1. 16.]
 [16.  9. 16.  4.  1.]
 [16.  1.  1. 16. 16.]
 [ 4.  9.  4. 16.  1.]
 [ 4.  1.  4. 16.  9.]]
Time of Squaring on GPU with gpuarray
0.058682s

```

```

In [3]: runfile('G:/cuda opencv book material/CUDA book code/Chapter11/gpu_dot.py',
              wdir='G:/cuda opencv book material/CUDA book code/Chapter11')
Answer of Dot Product using numpy
633.0
Time taken for Dot Product using numpy
0.03769350051879883 s
Answer of Dot Product on GPU
633.0
Time taken for Dot Product on GPU
0.000108s
The computed dot product is correct

```

```

In [3]: runfile('G:/cuda opencv book material/CUDA book code/Chapter11/gpu_dot.py',
              wdir='G:/cuda opencv book material/CUDA book code/Chapter11')
Answer of Dot Product using numpy
633.0
Time taken for Dot Product using numpy
0.03769350051879883 s
Answer of Dot Product on GPU
633.0
Time taken for Dot Product on GPU
0.000108s
The computed dot product is correct

```

$$\begin{bmatrix} 4 & 1 & 1 \\ 1 & 2 & 3 \\ 1 & 1 & 3 \end{bmatrix} + \begin{bmatrix} 3 & 3 & 4 \\ 2 & 3 & 3 \\ 4 & 1 & 1 \end{bmatrix} = \begin{bmatrix} 4 * 3 + 1 * 2 + 1 * 4 & 16 & 20 \\ 1 * 3 + 2 * 2 + 3 * 4 & 12 & 13 \\ 1 * 3 + 1 * 2 + 3 * 4 & 9 & 10 \end{bmatrix}$$

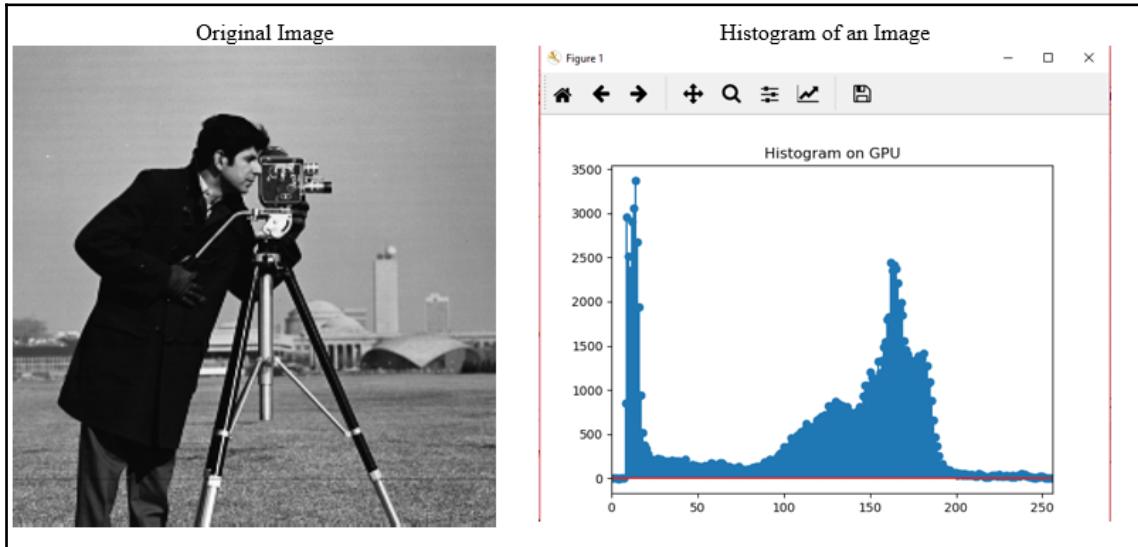
```
In [7]: runfile('G:/cuda opencv book material/CUDA book code/Chapter11/matrix_mmulfinal.py', wdir='G:/cuda opencv book material/CUDA book code/Chapter11')
*****
Matrix A:
[[4.  1.  1.]
 [1.  2.  3.]
 [1.  1.  3.]]
*****
Matrix B:
[[3.  3.  4.]
 [2.  3.  3.]
 [4.  1.  1.]]
*****
Matrix Multiplication result:
[[18. 16. 20.]
 [19. 12. 13.]
 [17.  9. 10.]]
The computed matrix multiplication is correct
```

```
In [4]: runfile('G:/cuda opencv book material/CUDA book code/Chapter12/element_wise_addition.py', wdir='G:/cuda opencv book material/CUDA book code/Chapter12')
Addition of 1000000 element of GPU
0.000629s
The sum computed on GPU is correct
```

```
In [6]: runfile('G:/cuda opencv book material/CUDA book code/Chapter12/gpu_dot.py',
wdir='G:/cuda opencv book material/CUDA book code/Chapter12')
Vector A
[0 1 2 3 4]
Vector B
[0 1 2 3 4]
The computed dot product using reduction:
30
Dot Product on GPU
2.546338s
```

```
In [3]: runfile('G:/cuda opencv book material/CUDA book code/Chapter12/gpu_scan.py',
wdir='G:/cuda opencv book material/CUDA book code/Chapter12')
The input data:
[7 5 9 2 9 7 7 7 2 6]
The computed cumulative sum using Scan:
[ 7 12 21 23 32 39 46 53 55 61]
Cumulative Sum on GPU
0.002032s
```

# Chapter 12: Basic Computer Vision Applications Using PyCUDA



```
In [4]: runfile('G:/cuda opencv book material/CUDA book code/Chapter12/histogram_without_shared.py', wdir='G:/cuda opencv book material/CUDA book code/Chapter12')
```

Time for Calculating Histogram on GPU without shared memory

0.001040s

Time for Calculating Histogram using OpenCV

0.001040s

```
In [5]: runfile('G:/cuda opencv book material/CUDA book code/Chapter12/histogram.py', wdir='G:/cuda opencv book material/CUDA book code/Chapter12')
```

Time for Calculating Histogram on GPU with shared memory

0.000839s

